

Knowledge Graphs for AI in the Enterprise

François Scharffe

Tech Innovations Forum

Columbia University Libraries

11/25/2019



About me

- François Scharffe
- @lechatpito on Twitter
- PhD from University of Innsbruck, in EU Semantic Web lab
- Senior lecturer at the School of Professional Studies in the Applied Analytics Master in Sciences, Columbia University
- Consulting through my company The Data Chefs
- Organizer of the Knowledge Graph Conference


The AI/ML/data science process

Define the business problem

What is the problem at hand?
Is there a question to answer?
Is it an open problem?

Build the dataset

What are the data sources needed to solve the problem? Which ones are available? How to integrate the data? How often will the analysis performed on the data?



Evaluate results

Run the model over the data. Evaluate the results. Should the problem be reformulated? Is the dataset sufficient to answer the questions?

Design the model

Will a set of queries able to answer our business problem? Do we need a statistical model? Which model is appropriate? What are the features? The parameters of the models?


The The AI/ML/data science process

Define the business problem

What is the problem at hand?
Is there a question to answer?
Is it an open problem?

Build the dataset

What are the data sources needed to solve the problem? Which ones are available? How to integrate the data? How often will the analysis performed on the data?




Evaluate results

Run the model over the data. Evaluate the results. Should the problem be reformulated? Is the dataset sufficient to answer the questions?

Design the analytical model

Will a set of queries able to answer our business problem? Do we need a statistical model? Which model is appropriate? What are the features? The parameters of the models?

Building a data set is time consuming



What data scientists spend the most time doing


- Building training sets: 3%
- Cleaning and organizing data: 60%
- Collecting data sets; 19%
- Mining data for patterns: 9%
- Refining algorithms: 4%
- Other: 5%

Source Crowdfunder http://visit.crowdfunder.com/rs/416-ZBE-142/images/CrowdFunder_DataScienceReport_2016.pdf

See also

<https://www.forbes.com/sites/gilpress/2016/03/23/data-preparation-most-time-consuming-least-enjoyable-data-science-task-survey-says/#54d5863f6f63>

Building a data set is time consuming




What data scientists spend the most time doing

- Cleaning and organizing data: 60%
- Collecting data sets; 19%
- Mining data for patterns: 7%
- Refining algorithms: 4%
- Other: 5%

Credits Crowdfunder http://visit.crowdfunder.com/rs/416-ZBE-142/images/CrowdFunder_DataScienceReport_2016.pdf

See also

<https://www.forbes.com/sites/gilpress/2016/03/23/data-preparation-most-time-consuming-least-enjoyable-data-science-task-survey-says/#54d5863f6f63>



That's 80% of
their time !!!

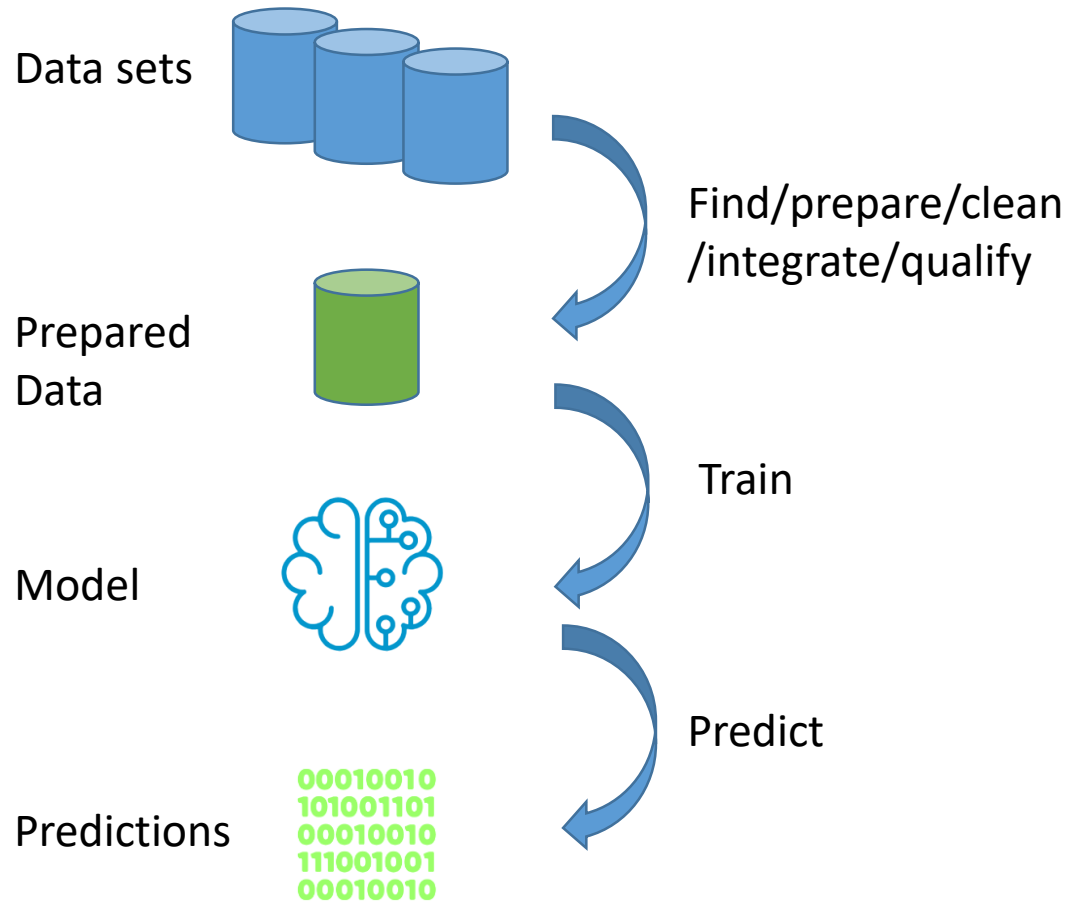
Knowledge Graphs come to the rescue

Many data sources are **integrated** so that they can be **reused** across applications and models


- The hard work is performed upstream
 - Data quality
 - Entity matching
 - Schema mapping
 - ETL
 - Ontology definition

KGs for smarter AI

- AI today



- AI tomorrow



Knowledge Graphs for AI

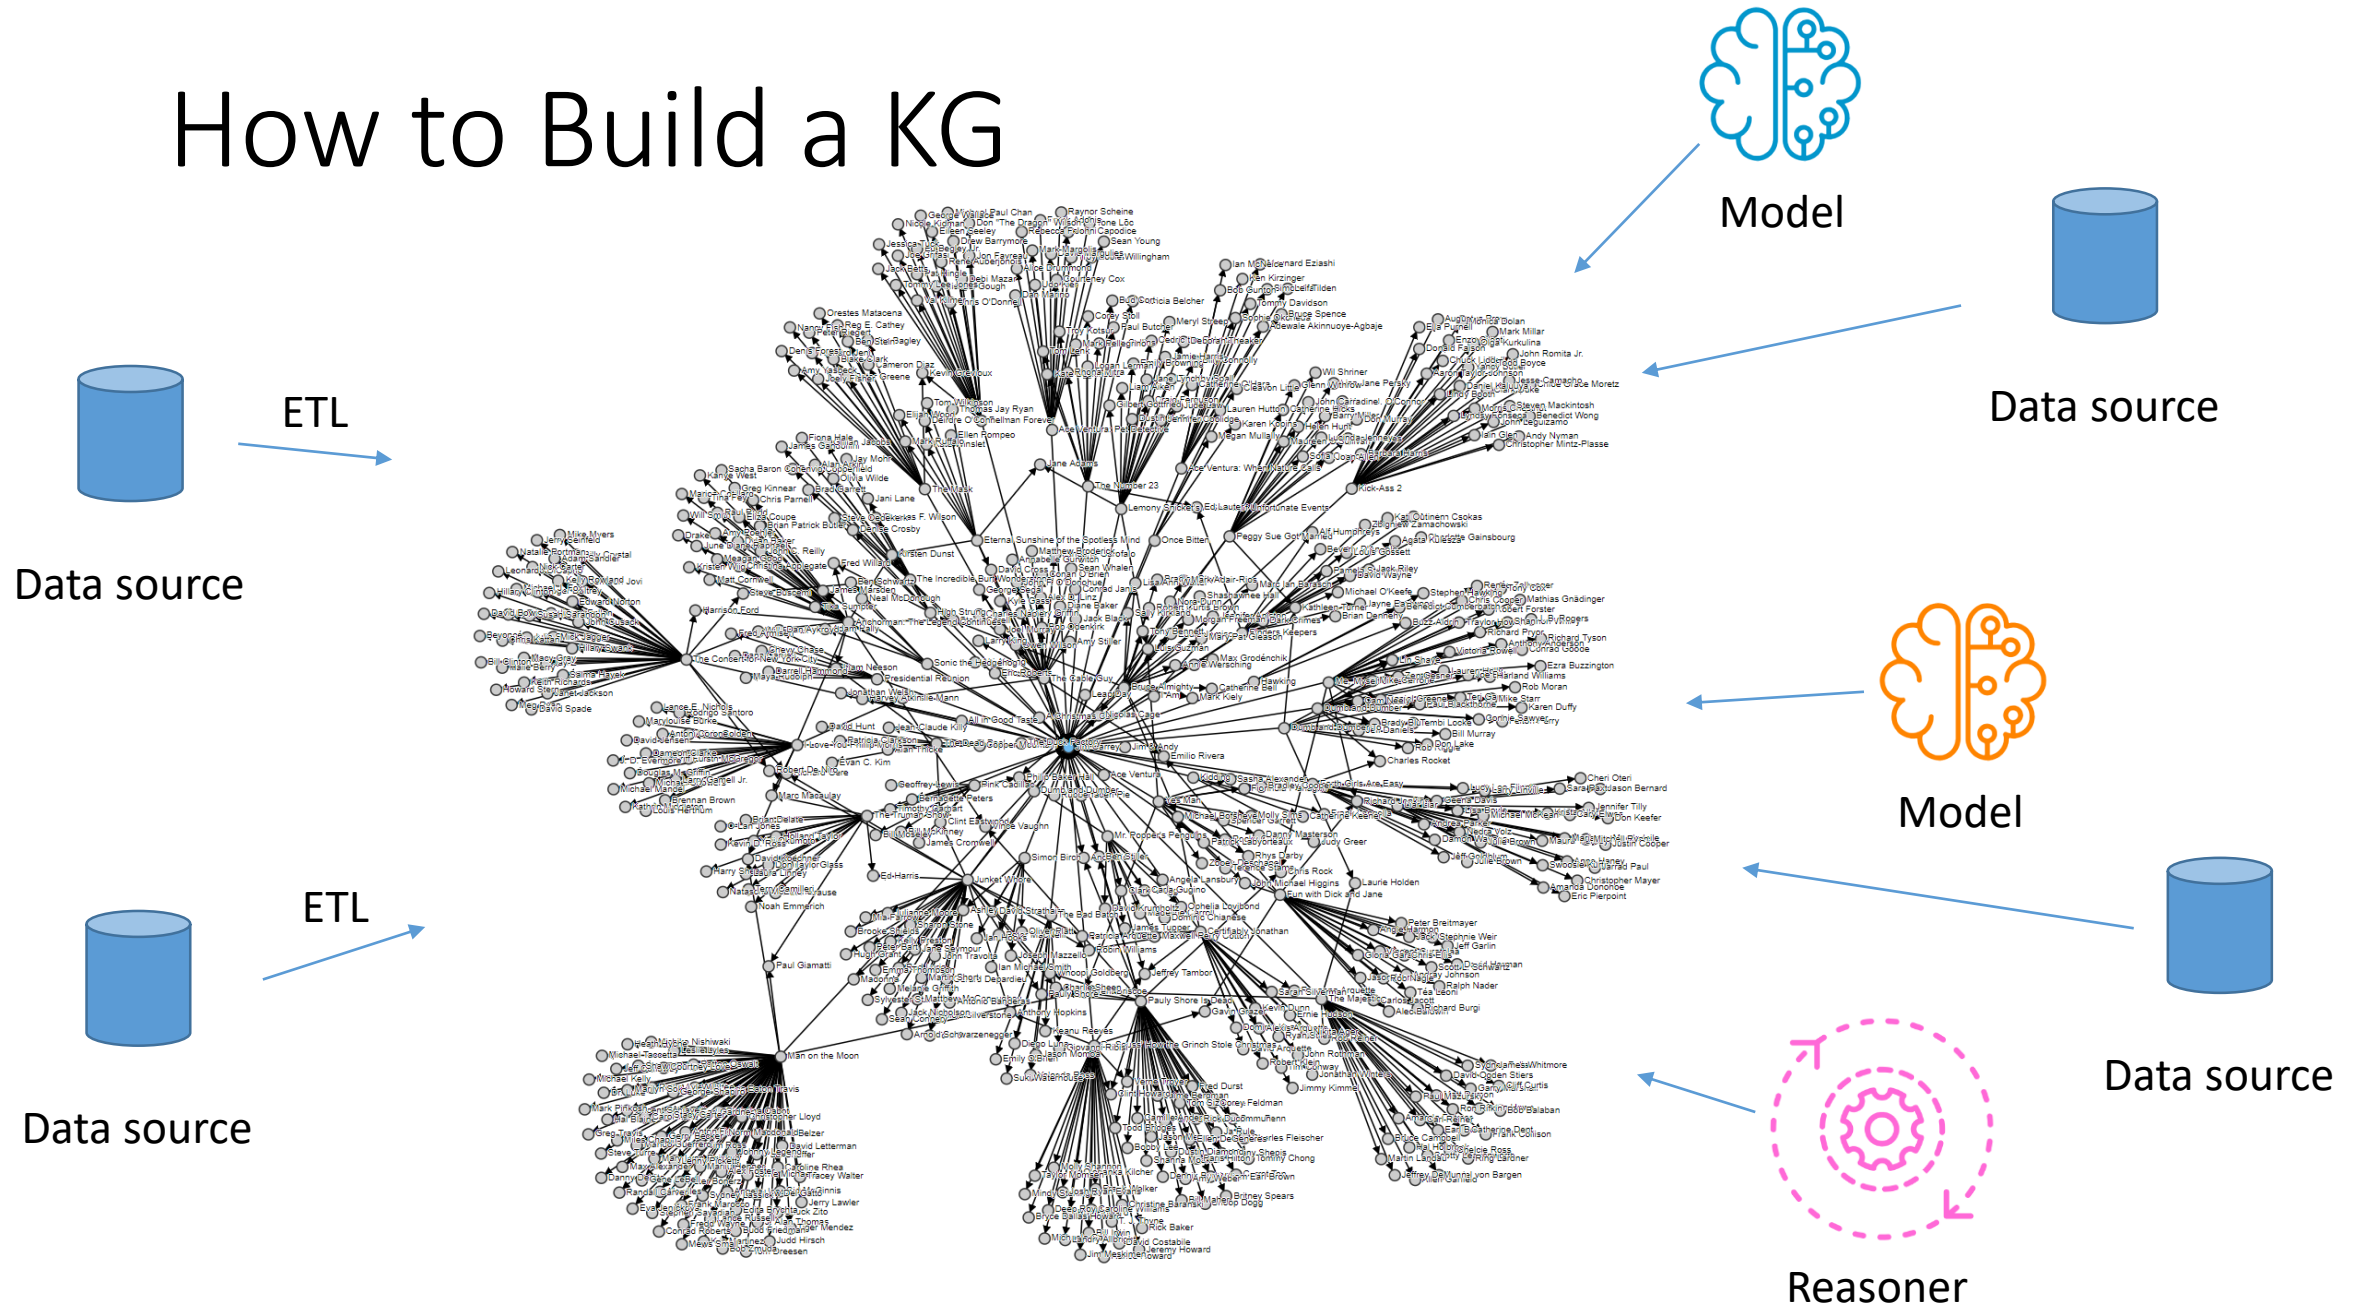
Intelligence needs a memory that
can be used to **reason logically**
about facts in the world

Knowledge graphs allows data to be
represented as facts

A brief history of knowledge graphs:

<http://knowledgegraph.today/paper.html>

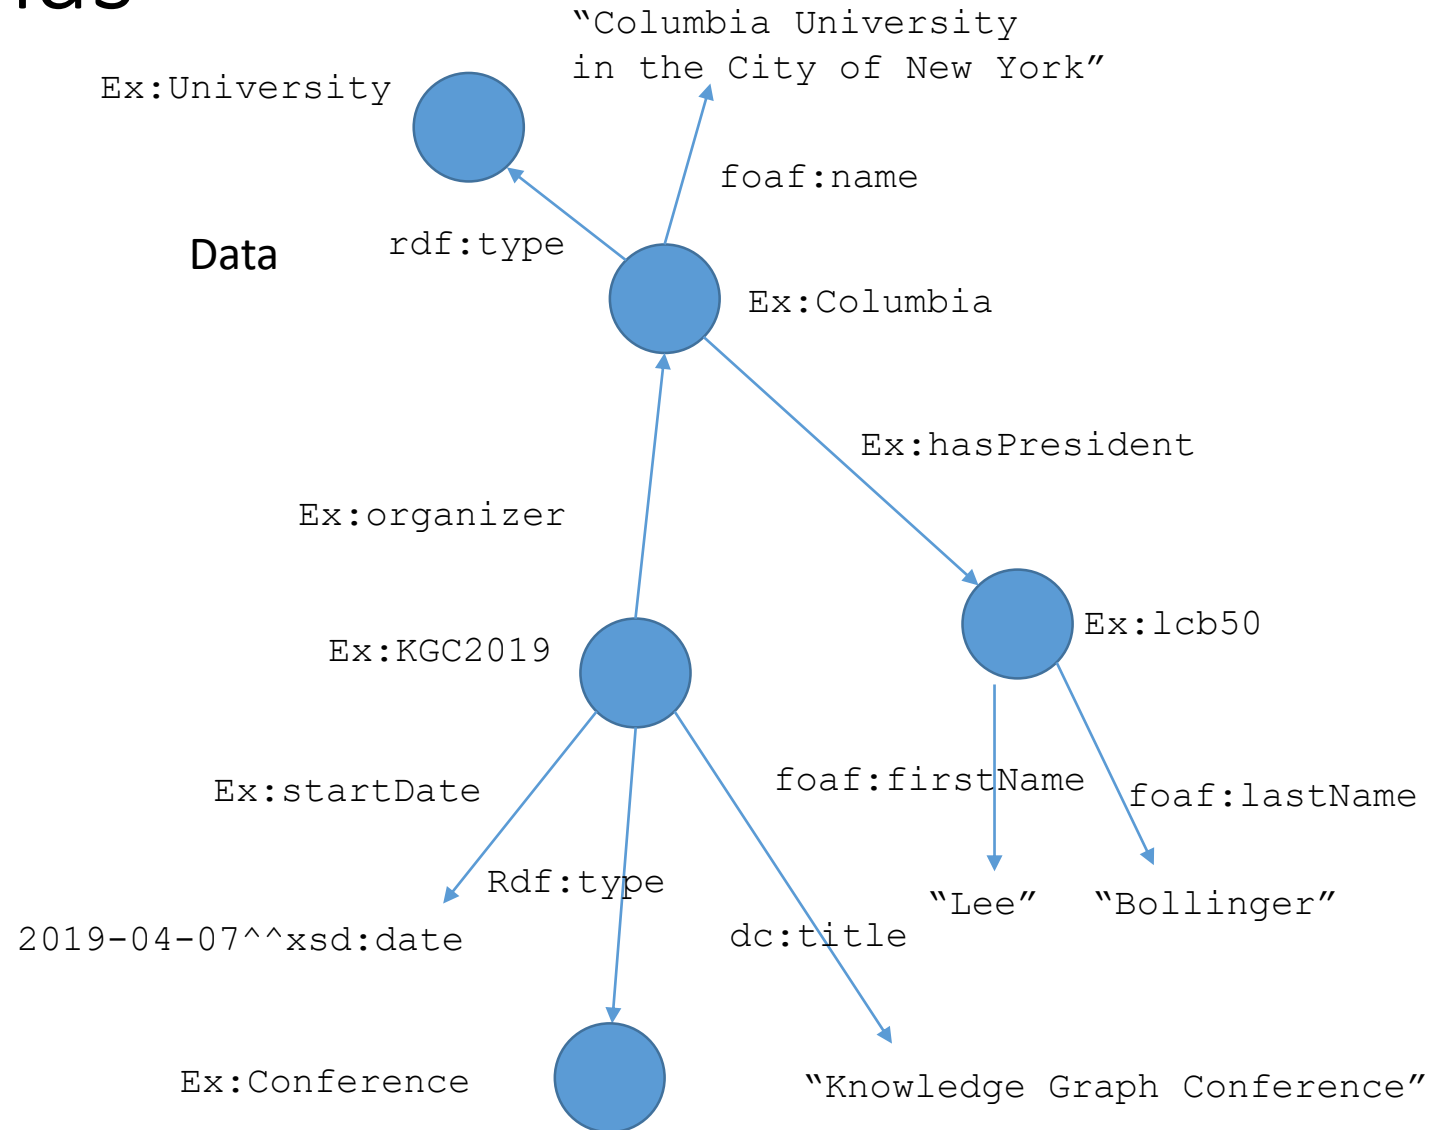
How to Build a KG



Graphs and Schemas

Ontology

```
Ex:University a rdfs:Class.  
Ex:organizer a rdfs:Property.  
Ex:organizer rdfs:domain Ex:Event.  
Ex:organizer rdfs:range foaf:Agent.  
Ex:Conference a rdfs:Class.  
Ex:Conference rdfs:subClassOf event:Event.  
rdfs:subClassOf foaf:Person.  
...
```



Graph Databases and Query Languages

- One open standard from the W3C
 - RDF the graph modeling language
 - SPARQL the graph query language
 - OWL the ontology language
- A mix of other more or less open languages
 - GSQL
 - GQL
 - Cypher
 - Gremlin
- Issues
 - fragmentation on one side
 - Complexity on the other. See [RDF spec](#)
- Ongoing [effort](#) to harmonize it all

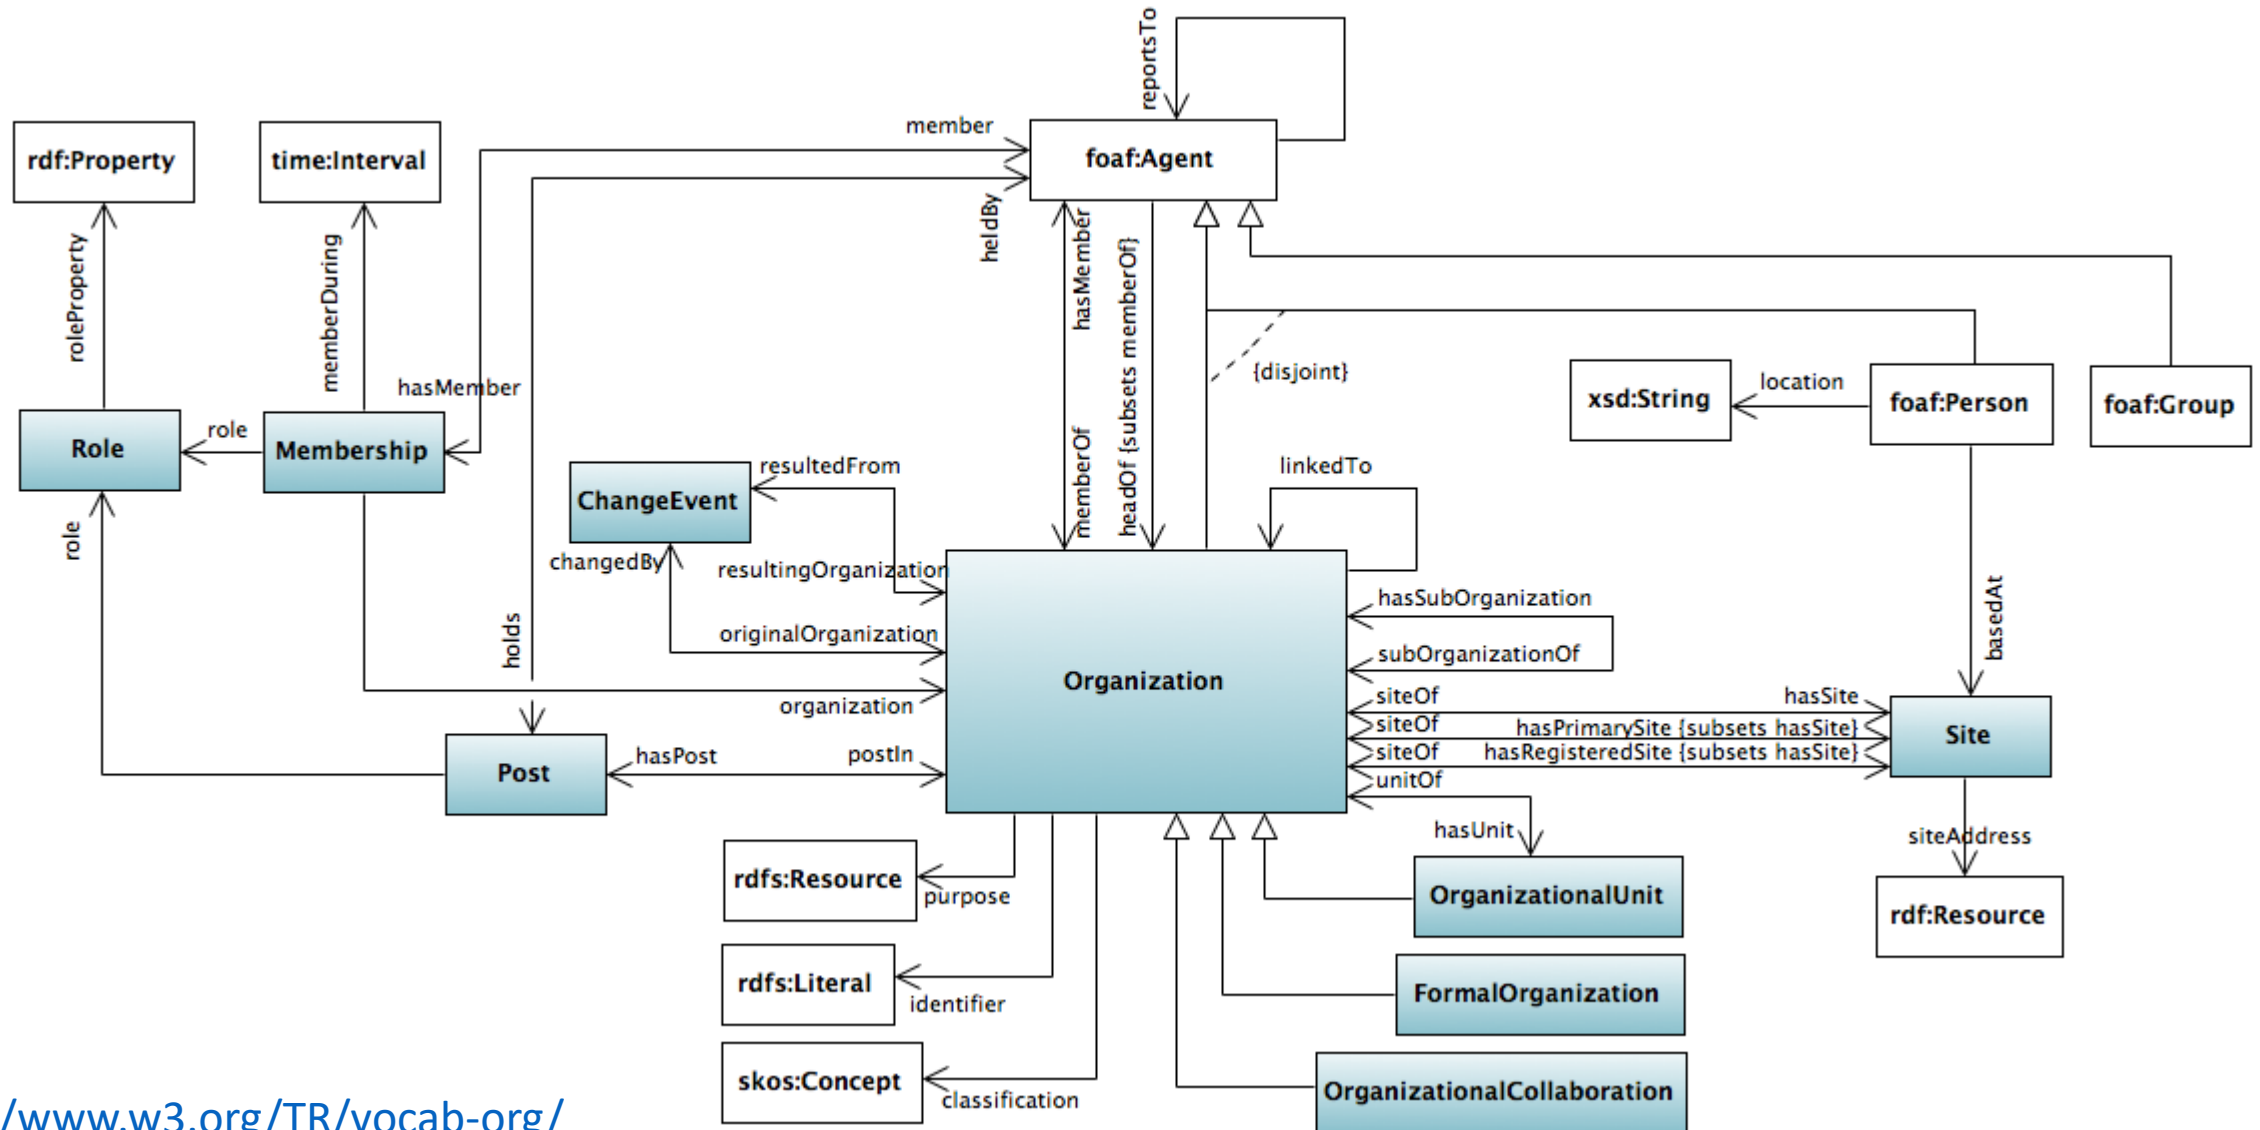


Ontolowhat?

An ontology is a **formal specification** of a **shared conceptualization** of a **domain**

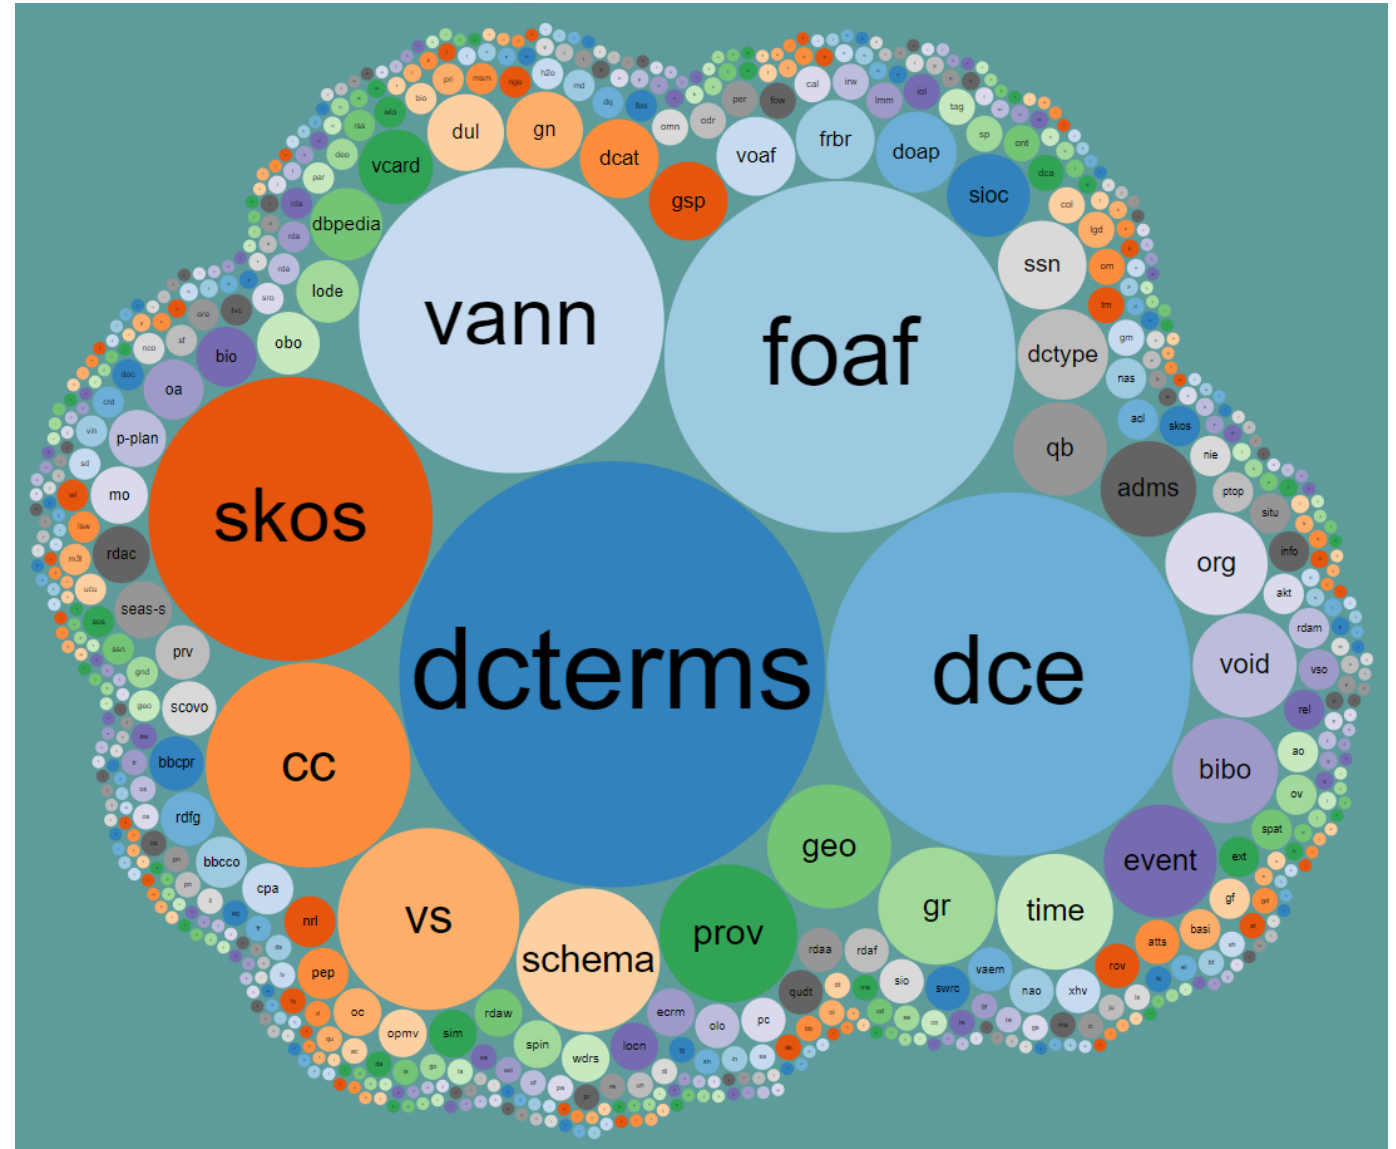
- a conceptualization: a description of the way we think about a domain
- a specification: there is a formal way of writing down the conceptualization
- formal: it is defined by axioms in a formal language
- Shared: it is meant to represent a community, it should be reusable across applications

Example: The Org Ontology



Ontologies

- <http://schema.org>
 - Open and collaborative development of an ontology for the annotation of web pages.
- [FIBO](#): “the semantic standard for the financial industry”
- Lots of small ontologies
<https://lov.linkeddata.es/dataset/lov/>

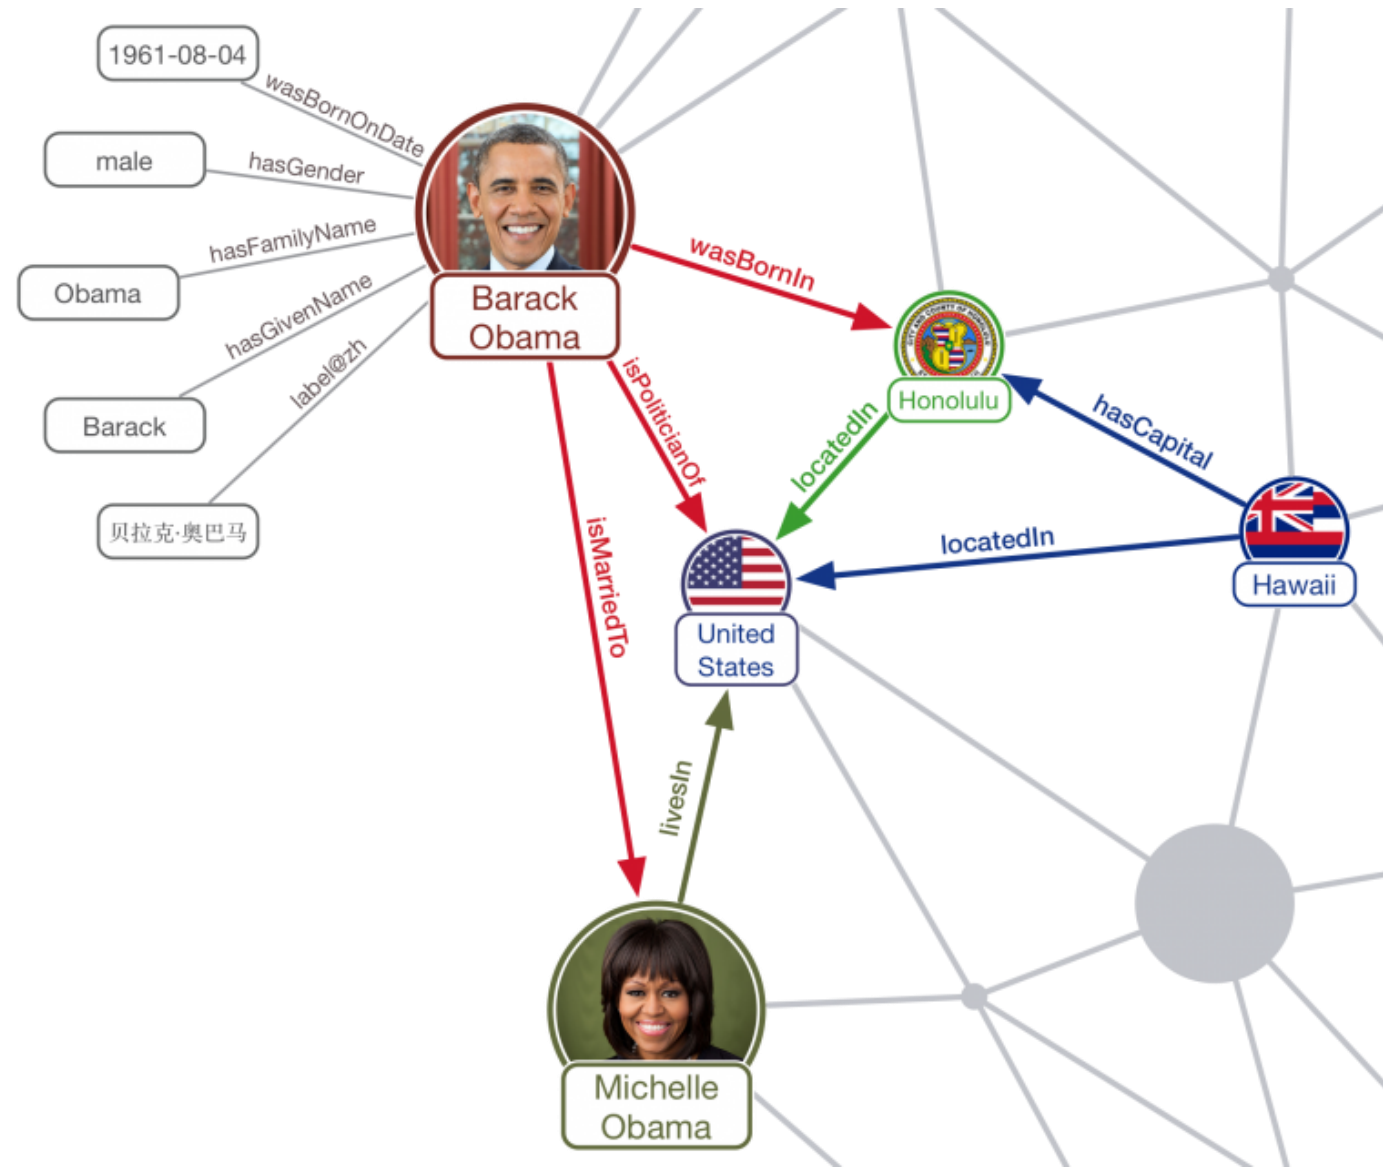


Graph vs Relational

- The inevitable question: why don't you use a SQL DB?
- There are many differences, let's cite the most relevant:
 - Graph DBs are more flexible
 - The data defines the graph structure. No a priori schema. Attributes (new links in the graph) can be added as you go and each resource has its own set of links.
 - Rich semantics are defined in ontologies
 - Ontologies define domains and are collaboratively created
 - This provide better integration with external datasets
 - Semantics can be tailored to the use case and the domain from very flexible to very crisp and constrained.
 - No need to join, path queries rule

Knowledge Graphs in the Wild


- Google knowledge graph
 - <https://developers.google.com/knowledge-graph/>
 - Powers Google search result cards
 - Open ontology available at <http://schema.org>
- Facebook OpenGraph
 - Now not open anymore!
 - Let you navigate and query for uses, likes, entities through an API
 - Powers Facebook AI for recommendations
- Microsoft Concept Graph
 - Concepts extracted from the Web
- Bing Satori
 - <https://azure.microsoft.com/en-us/services/cognitive-services/bing-entity-search-api/>
- Wikidata
 - Source of Wikipedia structured data
 - <http://www.Wikidata.org>
 - <https://angryloki.github.io/wikidata-graph-builder/>
- Refinitiv
 - <https://www.refinitiv.com/en/products/knowledge-graph-feed>
- And many more



The Knowledge Graphs Conference 2019



The Knowledge Graphs Conference 2020



<https://www.knowledgegraph.tech/>

<https://knowledgegraphconference.vhx.tv/>

[Groupe Slack](#)