

# **RDF Data Management & SPARQL Query Processing**

Cours développé par

**Federico Ulliana**

UM, LIRMM, INRIA GraphIK

Slides collected from Martin Theobald, Katja  
Hose, Ralf Schenkel, and Stratos Idreos

# NOSQL Umbrella

(RDF and SPARQL are the only standardized languages)

- Key-value databases are systems are about as simple as databases get, being in essence variations on the theme of a persistent hash table. Current examples include MemcacheDB, Tokyo Cabinet, Redis and SimpleDB.
- Document databases are key-value stores that treat stored values as semi-structured data instead of as opaque blobs. Prominent examples at the moment include CouchDB, MongoDB and Riak.
- Wide-column databases tend to draw inspiration from Google's BigTable model. Open-source examples include Cassandra, HBase and Hypertable.
- Graph databases include generic solutions like Neo4j, InfoGrid and HyperGraphDB as well as all the numerous RDF-centric solutions out there: AllegroGraph, 4store, Virtuoso, and many, many others.

# Objectifs du cours

*Comprendre le fonctionnement des systèmes d'interrogation de données RDF*

*Implémenter un mini-moteur d'évaluation de requêtes qui incorpore les idées vues en cours*

*Conduire des expériences permettant d'analyser les performances du système réalisé*

# Organisation

- 29 Octobre : RDF Stores 2CM + 1TP (début mini-projet)
- 19 Novembre: 1TP
- 26 Novembre: 2TP
- 3 Décembre: 2TP
- 10 Décembre: 2TP

# Plan

- **RDF and SPARQL**
- **RDF Row-stores**
- **RDF Column-stores**
- **RDF Graph-stores**

# RDF Triples

<Albert\_Einstein, isA, physicist>

<Albert\_Einstein, bornIn, Ulm>

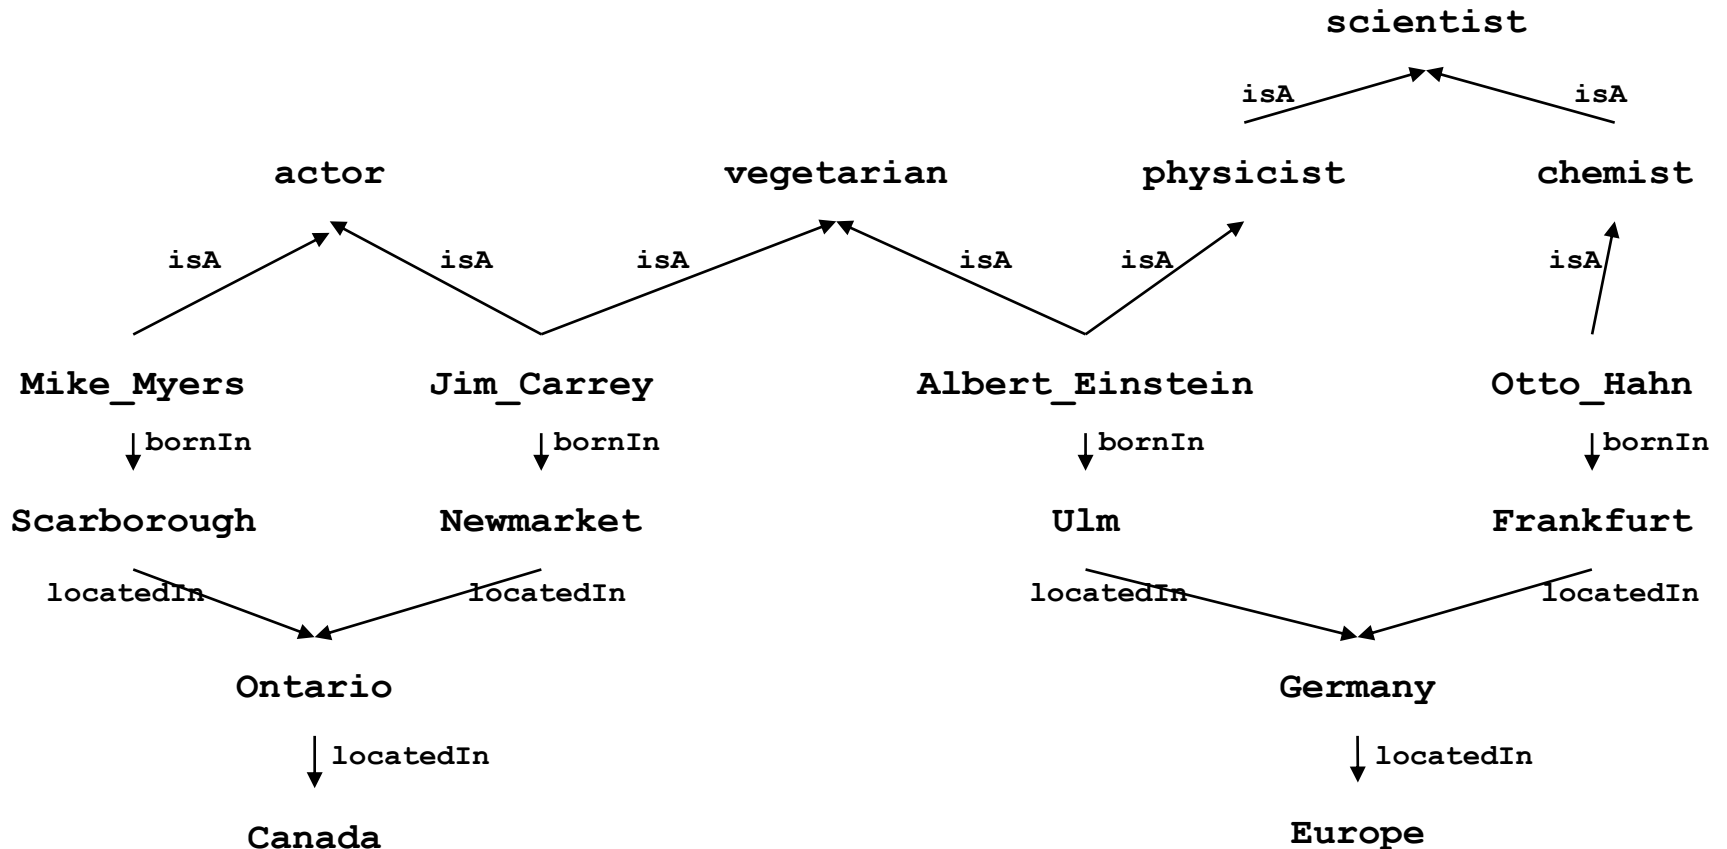
<Albert\_Einstein, isA, vegetarian> ...

# Graph Notation

<Albert\_Einstein, isA, physicist>


<Albert\_Einstein, bornIn, Ulm>

<Albert\_Einstein, isA, vegetarian> ...



# RDF is more than a Graph Database

<Albert\_Einstein, isA, physicist>

<isA, rdfs:subPropertyOf, rdfs:subClassOf> 

SELECT   ?x ?y

WHERE {

    ?x ?p ?y .

    ?p rdfs:subPropertyOf, rdf:subClassOf}

}



# RDF is more than a Graph Database

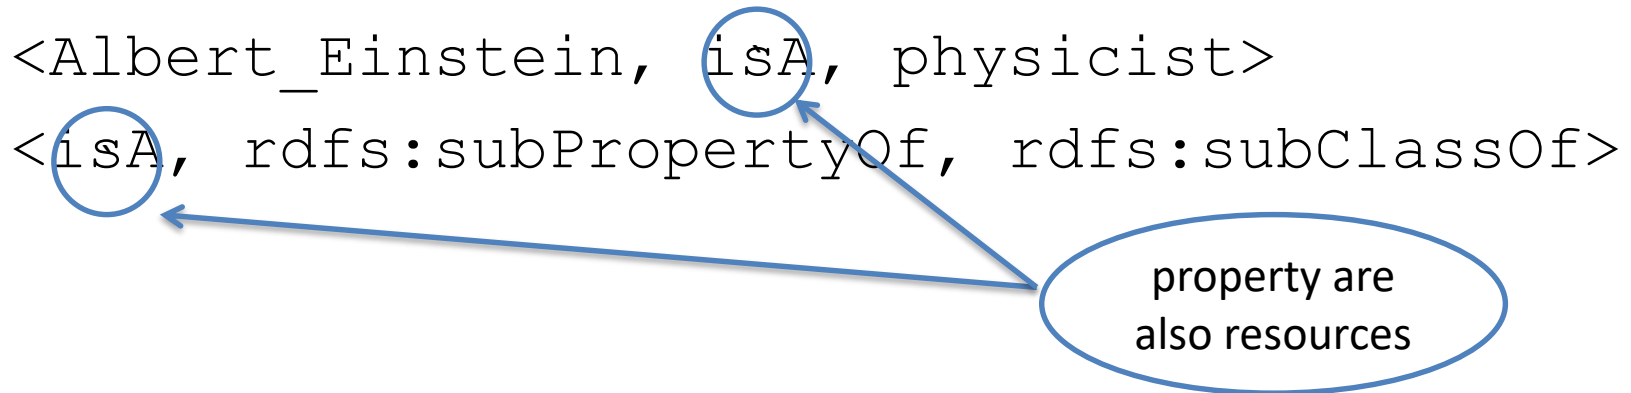
`<Albert_Einstein, isA, physicist>`  
`<isA, rdfs:subPropertyOf, rdfs:subClassOf>`

property are also reources

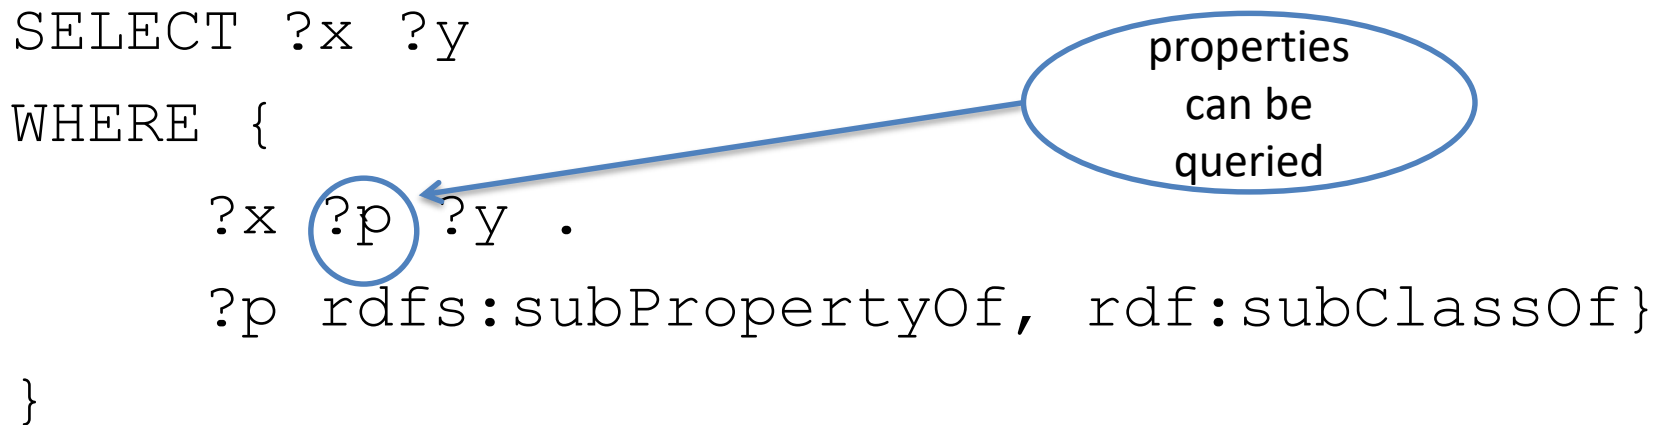
```
SELECT ?x ?y
WHERE {
    ?x ?p ?y .
    ?p rdfs:subPropertyOf, rdf:subClassOf}
}
```

# RDF is strictly more than a Graph Database

`<Albert_Einstein, isA, physicist>`  
`<isA, rdfs:subPropertyOf, rdfs:subClassOf>`




```
SELECT ?x ?y
WHERE {
  ?x ?p ?y .
  ?p rdfs:subPropertyOf, rdf:subClassOf}
}
```



# Why caring about RDF data?

## Open Data are gaining momentum !



More than 30 billion triples in more than 200 sources across the LOD cloud  
**DBpedia: 3.4 million entities, 1 billion triples**

# Queries can be complex, too

```
SELECT DISTINCT ?a ?b ?lat ?long WHERE
{ ?a dbpedia:spouse ?b.
  ?a dbpedia:wikilink dbpediares:actor.
  ?b dbpedia:wikilink dbpediares:actor.
  ?a dbpedia:placeOfBirth ?c.
  ?b dbpedia:placeOfBirth ?c2.
  ?c owl:sameAs ?c2.
  ?c2 pos:lat ?lat.
  ?c2 pos:long ?long.
}
```



# SPARQL 1.0 / 1.1

- Query language for RDF suggested by the W3C
- SPARQL main building block: **triple patterns**

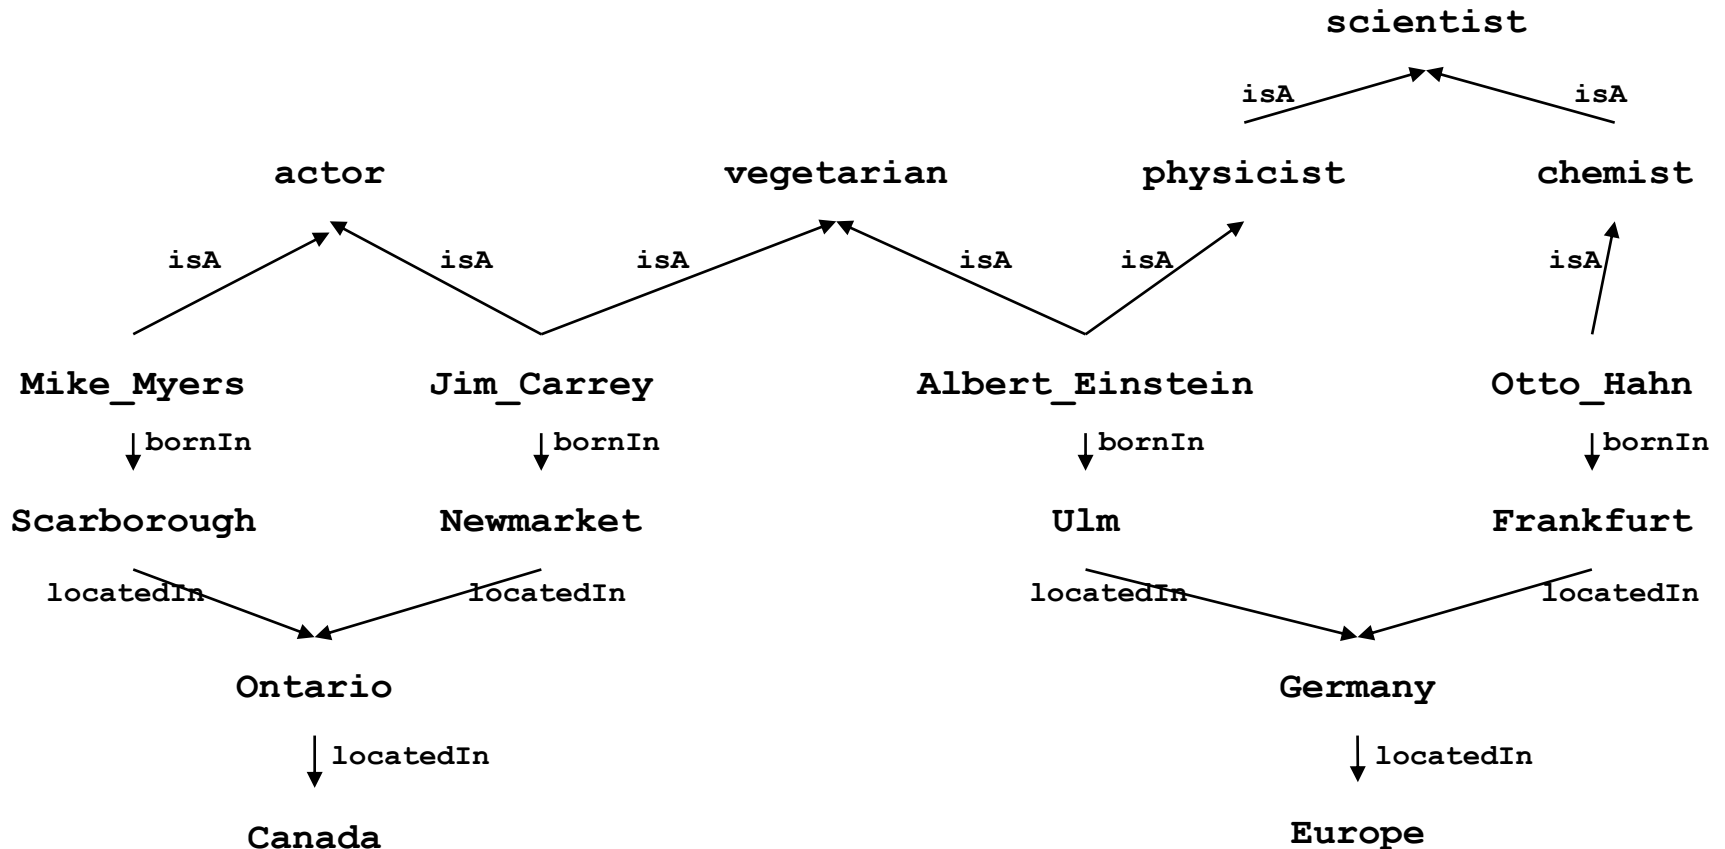
`?x wrote_book ?y`

- Like **select-project-join** for relational databases

# SPARQL – Example

Example query:

**Find all actors from Ontario** (that are in the knowledge base)



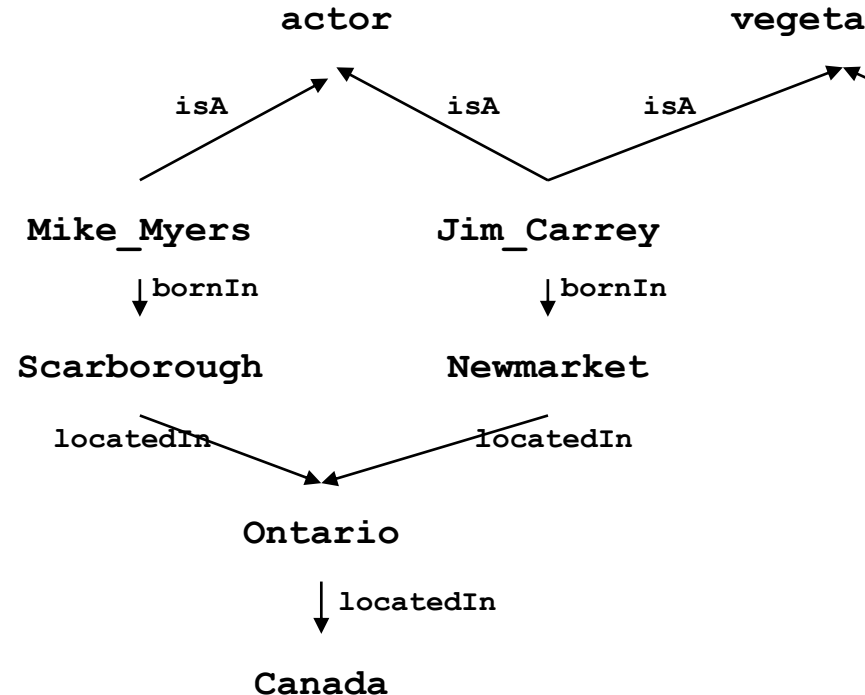

# SPARQL – Example

Example query:

**Find all actors from Ontario** (that are in the knowledge base)

```
SELECT ?person WHERE {  
  ?person isA actor.  
  ?person bornIn ?loc .  
  ?loc locatedIn Ontario .  
}
```

Find **subgraphs** of this form:



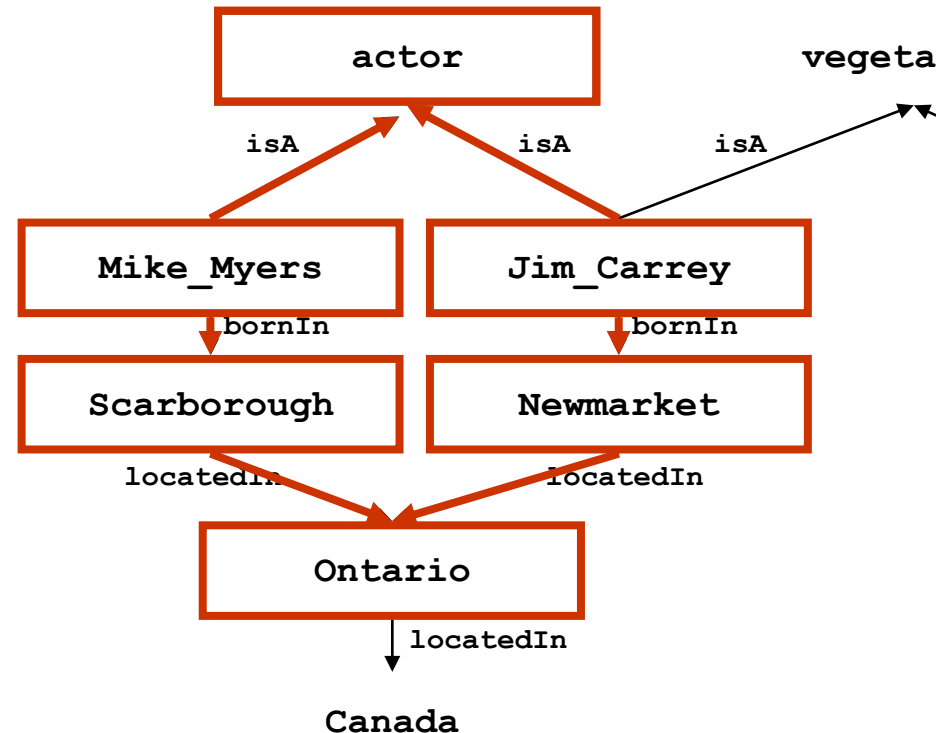

# SPARQL – Example

Example query:

**Find all actors from Ontario** (that are in the knowledge base)

```
SELECT ?person WHERE {  
  ?person isA actor.  
  ?person bornIn ?loc .  
  ?loc locatedIn Ontario .  
}
```

Find **subgraphs** of this form:





# Questions

- How to store RDF data ?
- How to query RDF data ?
- We will study three main approaches
  - Row-stores
  - Column-stores
  - Graph-stores

**ROW-STORES**

# Row-store

Classic relational database, storing relations by rows.  
(Postgres, Oracle, DB2, MySQL, ... )

<b>product</b>	<b>country</b>	<b>sales</b>
car	US	40K
bike	US	7K
	...	



**row1**


car@US@40K

**row2**

bike@US@7K

...

# RDF in a Relational Row-store



1 triple = 1 edge  
in the RDF Graph

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
	...	



**row1**

alice@work  
s\_for@EDF

**row2**

alice@live  
sIn@Paris

...


# Giant-Table (Jena, HexaStore, RDF-3X)

1. Store triples in one **giant** 3-attribute table
2. Convert SPARQL to equivalent SQL
3. *Magic* : the database will do the rest

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

**Giant**  
1 Billion  
Triples  
=  
1 Billion  
lines

# Conversion of SPARQL to SQL

triple pattern	→	FROM/WHERE
Shared variables	→	(self)JOIN conditions 
Constants	→	WHERE conditions
FILTER conditions	→	WHERE conditions
OPTIONAL clauses	→	OUTER JOINS
UNION clauses	→	UNION expressions

# Conversion of Triple Patterns (with Constants)

**SPARQL** >

```
SELECT ?x WHERE {?x lives_in ?y}
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# Conversion of Triple Patterns (with Constants)

**SPARQL** >

```
SELECT ?x WHERE {?x lives_in ?y}
```

**SQL** >

```
SELECT subject FROM Giant-Table  
WHERE predicate = "lives_in"
```

<b>subject</b>	<b>predicate</b>	<b>object</b>
alice	works_for	EDF
alice	lives_in	Paris
...		



# Conversion of Shared Variables

**SPARQL** >

```
SELECT * WHERE { ?x lives_in ?y.  
                  ?x works_for ?z }
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# Conversion of Shared Variables

**SPARQL** >

```
SELECT * WHERE { ?x lives_in ?y.  
                  ?x works_for ?z }
```

**SQL** >

```
SELECT L.subject, L.object, W.object  
FROM Giant-Table as L, Giant-Table as W  
WHERE L.predicate = "lives_in"  
AND   W.predicate = "works_for"  
AND   L.subject = W.subject
```



subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# Conversion of FILTER conditions

**SPARQL** >

```
SELECT ?x WHERE { ?x lives_in ?y .  
                  FILTER(?y!="New York") }
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# Conversion of FILTER conditions

**SPARQL** >

```
SELECT ?x WHERE { ?x lives_in ?y .  
                  FILTER(?y!="New York") }
```

**SQL** >

```
SELECT subject  
FROM Giant-Table  
WHERE predicate = "lives_in"  
      AND object != "New York"
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# Conversion of UNION clauses

**SPARQL** >

```
SELECT ?x WHERE {  
    { ?x lives_in ?y }  
    UNION { ?x works_for ?z } }
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# Conversion of UNION clauses

**SPARQL** >

```
SELECT ?x WHERE {  
    { ?x lives_in ?y }  
    UNION { ?x works_for ?z } }
```

**SQL** >

```
SELECT subject FROM Giant-Table  
WHERE predicate = "lives_in" )  
UNION  
SELECT subject FROM Giant-Table  
WHERE predicate = "works_for"
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# Conversion of OPTIONAL clauses

**SPARQL** >

```
SELECT ?x WHERE { ?x lives_in ?y .  
                  OPTIONAL { ?x works_for ?z} }
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# Conversion of OPTIONAL clauses

**SPARQL** >

```
SELECT ?x WHERE { ?x lives_in ?y .  
                  OPTIONAL { ?x works_for ?z } }
```

**SQL** >

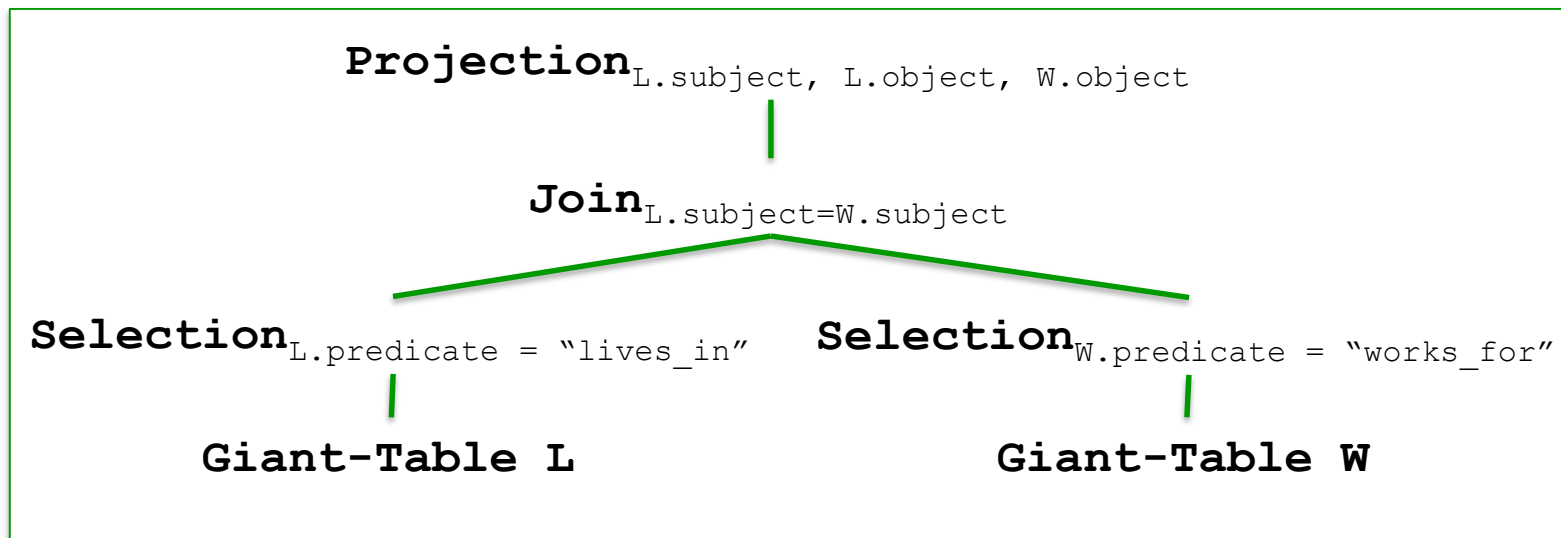
```
( SELECT      subject      FROM Giant-Table  
  WHERE      predicate = "lives_in" ) L  
LEFT OUTER JOIN  
  ( SELECT      subject      FROM Giant-Table  
    WHERE      predicate = "works_for" ) W  
ON (L.subject = W.subject)
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

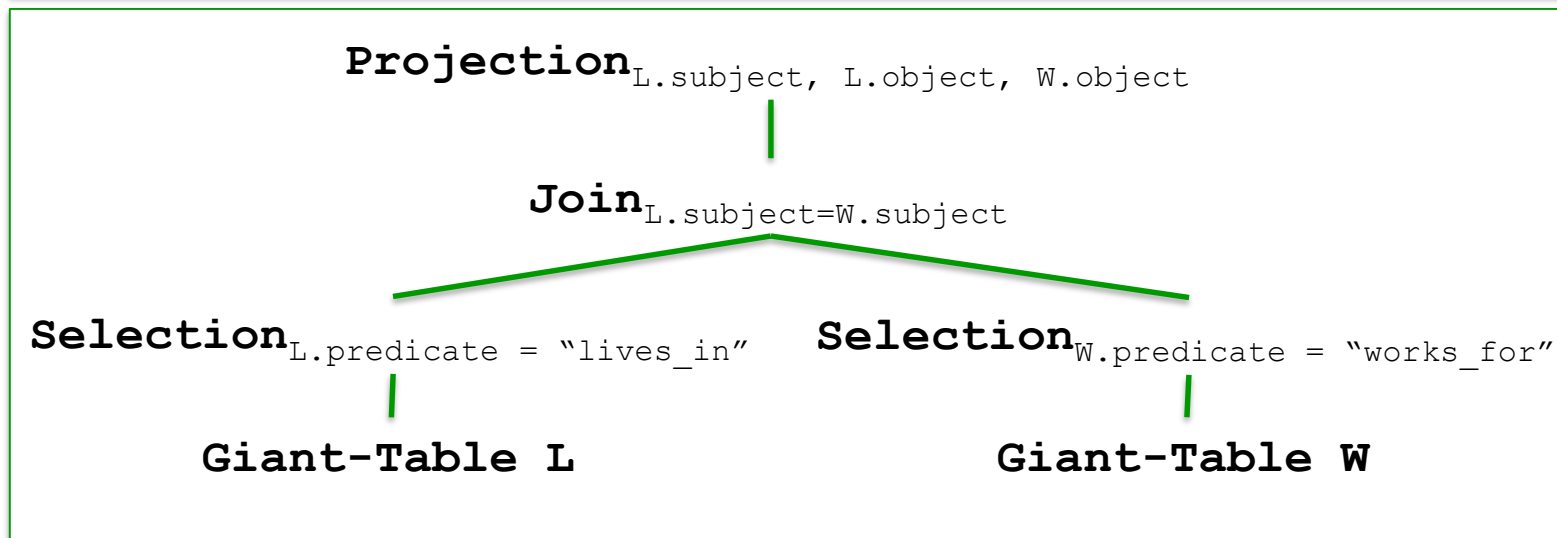
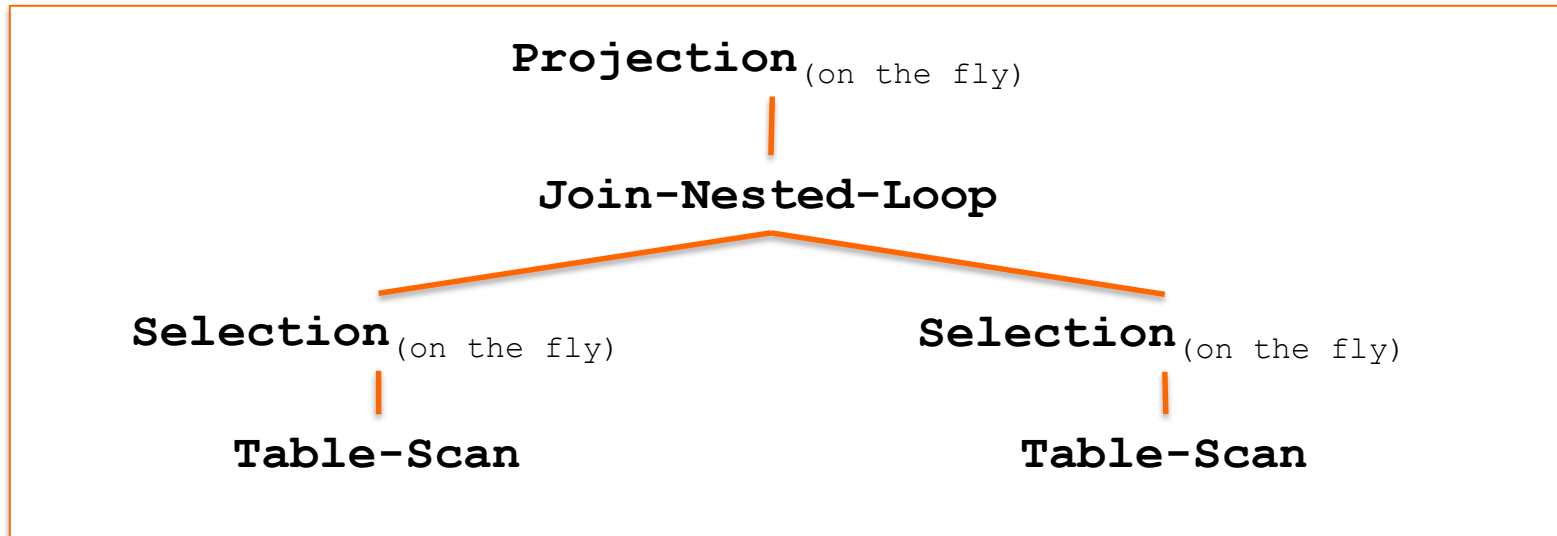


# From SQL to Relational Algebra

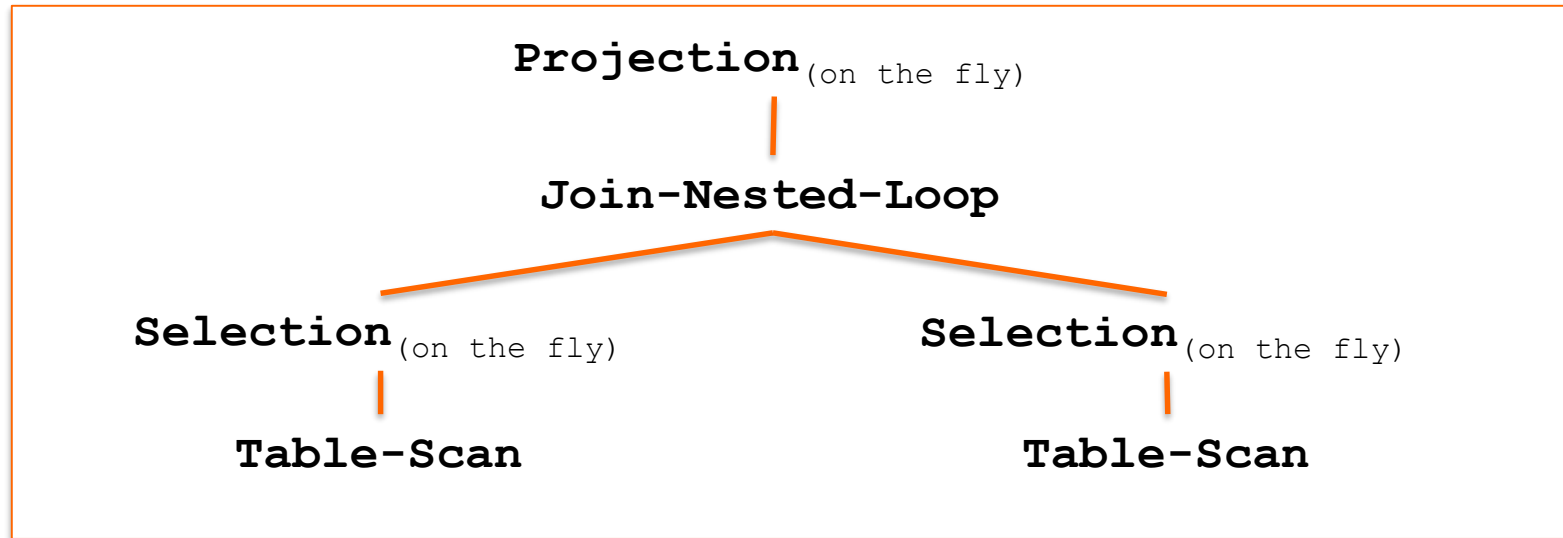
```
SELECT L.subject, L.object, W.object
FROM Giant-Table as L, Giant-Table as W
WHERE L.predicate = "lives_in"
AND    W.predicate = "works_for"
AND    L.subject = W.subject
```



# From Relational Algebra to Physical Plan



# From Relational Algebra to Physical Plan



- Now, performance depend on the row-store.



**row1**

alice@work  
s\_for@EDF

**row2**

alice@live  
sIn@Paris

...

Is that all?

# Is that all?

**Well, no.**

- Which other logical **schemas** can we use ?
- Which **indexes** should be built?  
(to support efficient evaluation of triple patterns)
- How can we **reduce storage space**?
- How can we find the **best execution plan**?

# Is that all?

**Well, no.**

- Which other logical **schemas** can we use ?
- Which **indexes** should be built?  
(to support efficient evaluation of triple patterns)
- How can we **reduce storage space**?
- How can we find the **best execution plan**?

**Existing databases need modifications:**

- flexible, extensible, generic storage not needed here
- cannot deal with multiple self-joins of a single table
- often generate bad execution plans

# **EXPLORING ALTERNATIVE RELATIONAL SCHEMAS**

# Problems with **Giant-Table**

- Too many joins, over a too large table.
- Alternative = many tables instead of one
- Property-Tables
- Clustered Property-Tables
- Property-Class



# Property-Tables

- A relational table for each single RDF property.

Giant-Table	subject	predicate	object
	alice	works_for	EDF
	alice	lives_in	Paris
	...		

**works\_for**

subject	object
alice	EDF
...	

**lives\_in**

subject	object
alice	Paris
...	

# The former conversion ...

**SPARQL** >

```
SELECT * WHERE { ?x lives_in ?y.  
                  ?x works_for ?z }
```

**SQL** >

```
SELECT L.subject, L.object, W.object  
FROM Giant-Table as L, Giant-Table as W  
WHERE L.predicate = "lives_in"  
AND   W.predicate = "works_for"  
AND   L.subject = W.subject
```

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris
...		

# ... now goes as follows

**SPARQL** >

```
SELECT * WHERE { ?x lives_in ?y.  
                  ?x works_for ?z }
```

**works\_for**

subject	object
alice	EDF
...	

**lives\_in**

subject	object
alice	Paris
...	

# ... now goes as follows

**SPARQL** >

```
SELECT * WHERE { ?x lives_in ?y.  
                  ?x works_for ?z }
```

**SQL** >

```
SELECT L.subject, L.object, W.object  
FROM lives_in as L, works_for as W  
WHERE L.subject = W.subject
```

**works\_for**

subject	object
alice	EDF
...	

**lives\_in**

subject	object
alice	Paris
...	

# Property-Tables

- Syntactically, we get smaller WHERE conditions
- Operationally, we avoid to self join a huge table
  - Keeping intermediary join result small is the key for efficiency in any database system

`works_for`

<b>subject</b>	<b>object</b>
alice	EDF
...	

`lives_in`

<b>subject</b>	<b>object</b>
alice	Paris
...	

# But properties can be correlated

YAGO: A Large Ontology from Wikipedia and WordNet

Group 1	Group 2	Group 3	Group 4	Group 5	Group 6
hasWebsite	isLocatedIn	hasGender	isCitizenOf	owns	created
isLocatedIn	isConnectedTo	isAffiliatedTo	wasBornIn	created	directed
owns	type	playsFor	livesIn	participatedIn	acted
created		wasBornIn		locatedIn	
	subClassOf		diedIn		influences

## **Examples**

<i>BMW</i>	<i>Los Angeles Airport</i>	<i>Alex Ferguson</i>	<i>John Belushi</i>	<i>Apple INC</i>	<i>Charlie Chaplin</i>
------------	--------------------------------	----------------------	---------------------	------------------	----------------------------

# Clustered Properties

`sport_man_property_cluster`

<b>subject</b>	<b>isAffiliatedTo</b>	<b>playsFor</b>	<b>hasGender</b>	<b>wasBornIn</b>
Ferguson	Manchester	Manchester	Male	Scotland
Ronaldo	UNICEF	Inter	Male	Brazil
...				

`employee_property_cluster`

<b>subject</b>	<b>works_for</b>	<b>lives_in</b>
alice	EDF	Paris
...		

# Recall the previous conversion ...

**SPARQL** >

```
SELECT * WHERE { ?x lives_in ?y.  
                  ?x works_for ?z }
```

**SQL** >

```
SELECT L.subject, L.object, W.object  
FROM lives_in as L, works_for as W  
AND L.subject = W.subject
```

**works\_for**

subject	object
alice	EDF
alice	Paris
...	

**lives\_in**

subject	object
alice	EDF
alice	Paris
...	



# now it goes as follows

**SPARQL** >

```
SELECT * WHERE {  
    ?x lives_in ?y.  
    ?x works_for ?z }  
}
```

**SQL** >

```
SELECT *  
FROM employee_property_cluster  
WHERE lives_in IS NOT NULL  
AND works_for IS NOT NULL
```

subject	works_for	lives_in
alice	EDF	Paris
...		

# Clustered Property-Tables

- Syntactically, we get even less join conditions
- Operationally, we avoid even more joins when properties are within a cluster
  - but we may still have to join two clusters!!



`works_for_lives_in_cluster`

<b>subject</b>	<b>works_for</b>	<b>lives_in</b>
alice	EDF	Paris
	...	

# Correlations do not always hold


<code>&lt;http://yago-knowledge.org/resource/isAffiliatedTo&gt;</code>	<b>2.635.440</b>
<code>&lt;http://yago-knowledge.org/resource/playsFor&gt;</code>	<b>2.575.219</b>
<code>&lt;http://yago-knowledge.org/resource/hasGender&gt;</code>	<b>345.794</b>
<code>&lt;http://yago-knowledge.org/resource/wasBornIn&gt;</code>	<b>172.541</b>

- Only 172.541 resources have a value for all properties.
- Clustering properties may waste a lot of storage (**nulls**)

# Correlations do not always hold

<http://yago-knowledge.org/resource/isAffiliatedTo>	2.635.440
<http://yago-knowledge.org/resource/playsFor>	2.575.219
<http://yago-knowledge.org/resource/hasGender	345.794
<http://yago-knowledge.org/resource/wasBornIn>	172.541

## many\_properties\_cluster

subject	isAffiliatedTo	playsFor	hasGender	wasBornIn
				null
		null		null
		null		null
		null		null
	null		null	null

# Attributes can have multiple values

many\_properties\_cluster

subject	isAffiliatedTo	playsFor	hasGender	wasBornIn
Ferguson	Manchester	Manchester	Male	Scotland
Ronaldo	UN	Inter	Male	Brazil
		...		

# Attributes can have multiple values

many\_properties\_cluster

subject	isAffiliatedTo	playsFor	hasGender	wasBornIn
Ferguson	Manchester	Manchester	Male	Scotland
Ronaldo	UNICEF	Inter	Male	Brazil
Ronaldo	<b>UNICEF</b>	Milan	<b>Male</b>	<b>Brazil</b>
Ronaldo	<b>UNICEF</b>	Barcelona	<b>Male</b>	<b>Brazil</b>
Ronaldo	<b>UNICEF</b>	RealMadrid	<b>Male</b>	<b>Brazil</b>
Ronaldo	<b>UNICEF</b>	Flamenco	<b>Male</b>	<b>Brazil</b>
		...		

- Note : by the way, the 4<sup>th</sup> normal form has been introduced exactly to avoid this.

# Leftover Triples

- Clustered Property Tables induce leftover triples
  - with none of the properties in a cluster
  - belonging to no class
  - extra joins between leftover-triples and clusters

leftover-triples

<b>subject</b>	<b>predicate</b>	<b>object</b>
alice	born_in	NY
EDF	located_in	Paris
...		

# The clustered-property table dilemma

- They are complex to design
  - If narrow: reduces nulls, increases unions/joins
  - If wide: reduces unions/joins, increases nulls



# Class-Property Tables

- A table contains all properties of the instances of a given class
- Has all inconveniences of the former method

`class:Book`

<b>subject</b>	<b>title</b>	<b>author</b>	<b>year</b>
ID1	XYZ	Joe Fox	2001
ID3	MNP	<b>null</b>	<b>null</b>
ID6	<b>null</b>	<b>null</b>	2004

# Property Tables: Pros and Cons

## Advantages:




- More in the spirit of existing relational systems
- Saves many self-joins over triple tables

## Disadvantages (mostly for clusters) :

- Potentially many NULL values
- Multi-value attributes problematic
- Schema changes very expensive

# **HEXASTORES : INDEXING IN RDF-3X**

# Hexastores

- RDF Systems introducing 6 indexing on triples
-  SPO,  PSO,  OSP
  - to access data by subject, property, or object
- SOP, POS, OPS
  - to cover all permutations


# Indexes

- Indexes are data-structures that allow a fast ( $\sim$ constant time) access to the stored data
- One can simply add them to boost the performance of any relational schemas
  - warning: an index can be larger than a database!
    - the real question is : when to stop indexing ??
- We look at a more original approach (RDF3X) that completely eliminates the schema design.

## Preprocessing: build a Dictionary for Strings


Map strings to unique integers (e.g., via hashing)

- Regular size (4-8 bytes), much easier to handle
- Dictionary usually kept in main memory

<code>http://example.fr/Alice</code>	→	1960	
<code>http://example.fr/Bob</code>	→	3795	
<code>http://example.fr/Charles</code>	→	4634	

**If not build carefully, this dictionary may break the original lexicographic sorting order  
⇒ FILTER conditions may be more expensive!**

# Dictionary




dictionary

id	string
1	alice
2	works_for
3	lives_in
4	EDF
5	Paris

Giant-Table

subject	pred.	object
1		

# Dictionary



dictionary


id	string
1	alice
2	works_for
3	lives_in
4	EDF
5	Paris

Giant-Table

subject	pred.	object
1	2	



# Dictionary




dictionary

id	string
1	alice
2	works_for
3	lives_in
4	EDF
5	Paris

Giant-Table

subject	pred.	object
1	2	4

# Dictionary




dictionary

id	string
1	alice
2	works_for
3	lives_in
4	EDF
5	Paris

Giant-Table

subject	pred.	object
1	2	4
1		

# Dictionary




dictionary

id	string
1	alice
2	works_for
3	lives_in
4	EDF
5	Paris

Giant-Table

subject	pred.	object
1	2	4
1	3	

# Dictionary



dictionary 

id	string
1	alice
2	works_for
3	lives_in
4	EDF
5	Paris

Giant-Table

subject	pred.	object
1	2	4
1	3	5



up to 10x  
faster!!

# RDF3X Storage and Indexing

- Giant-Table model + ad-hoc implementation (no RDBMs as we have seen before)
- Ad-hoc implementation here means that the Giant-Table is actually fused with indexes (we will see this next)

# What triple patterns are found in queries ?

( s    p    o )

( s    p    **?x** )

( s    **?x**    o    )

( **?x**    p    o    )


( s    **?x**    **?y** )

( **?x**    p    **?y** )

( **?x**    **?y**    o    )

( **?x**    **?y**    **?z** )


# So we can store triples pattern-wise



Giant-Table<SPO>

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris

# So we can store triples pattern-wise




Giant-Table<SOP>

subject	object	predicate
alice	EDF	works_for
alice	Paris	lives_in




# So we can store triples pattern-wise



Giant-Table<OPS>

object	predicate	subject
EDF	works_for	alice
Paris	lives_in	alice


# So we can store triples pattern-wise



Giant-Table<POS>

predicate	object	subject
works_for	EDF	alice
lives_in	Paris	alice

# So we can store triples pattern-wise



Giant-Table<PSO>

<b>predicate</b>	<b>subject</b>	<b>object</b>
works_for	alice	EDF
lives_in	alice	Paris

# Why ? Because we deal with row-stores

Giant-Table<SPO>

subject	predicate	object
alice	works_for	EDF
alice	lives_in	Paris

- Easier to match (s p ?x) patterns if stored as



<b>row1</b>	alice@work s_for@EDF	<b>row2</b>	alice@live sIn@Paris	...
-------------	-------------------------	-------------	-------------------------	-----

# Why ? Because we deal with row-stores

Giant-Table<POS>

<b>predicate</b>	<b>object</b>	<b>subject</b>
works_for	EDF	alice
lives_in	Paris	alice

- Easier to match (**?x** p o) patterns if stored as

up to 3x  
faster!!




**row1** works\_for@  
EDF@alice

**row2** livesIn@Pa  
ris@alice

...


# RDF3X Storage and Indexing

- Similarly RDF3X create 6 indexes
  - SPO ; SOP ; PSO ; POS ; OSP ; OPS
- SPO (Intuitively)

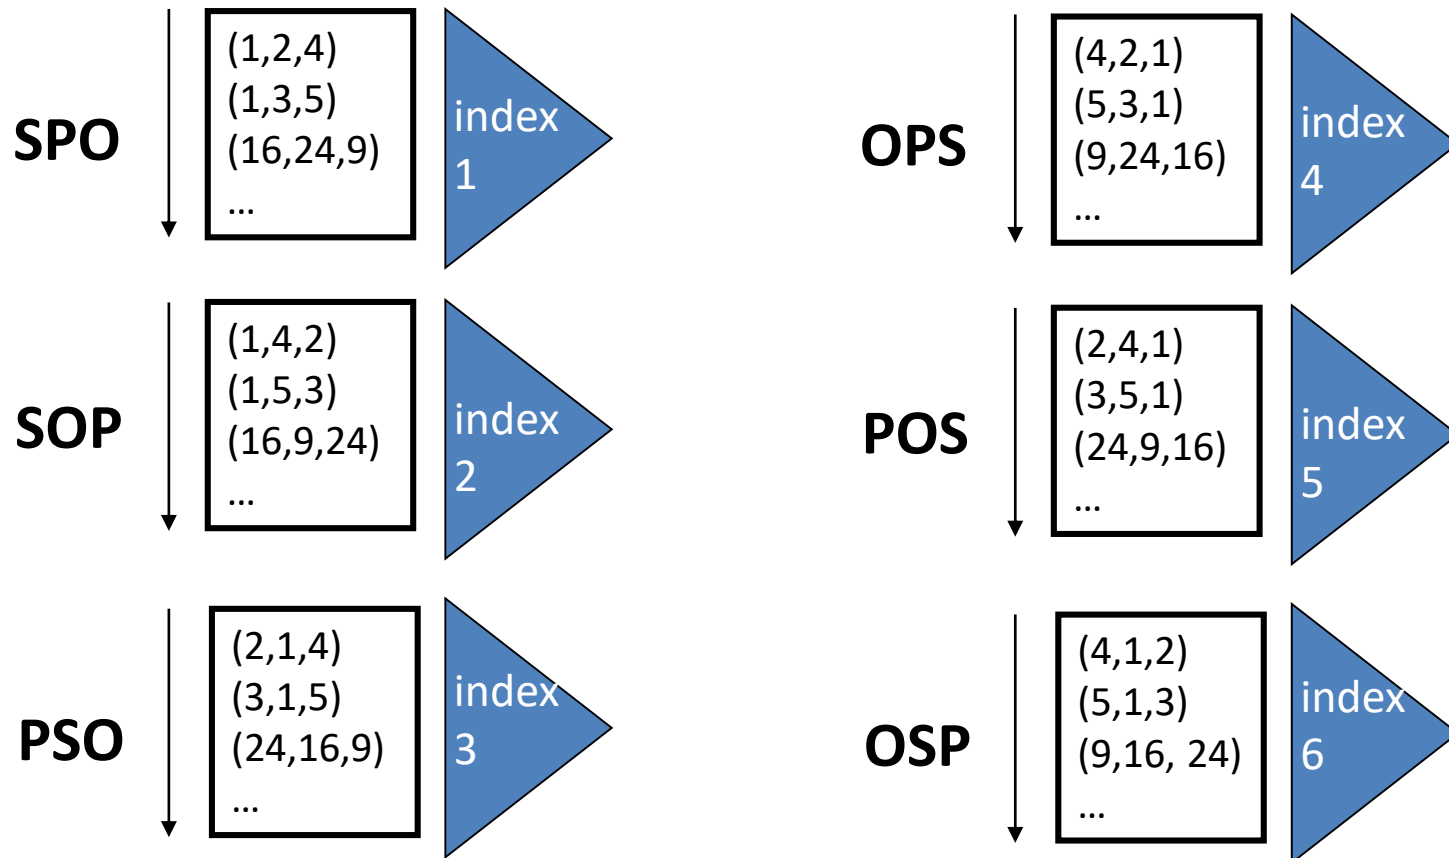


# RDF3X Storage and Indexing

- Similarly RDF3X create 6 indexes
  - SPO ; SOP ; PSO ; POS ; OSP ; OPS
- SPO (in reality) : (B+)-trees over triples



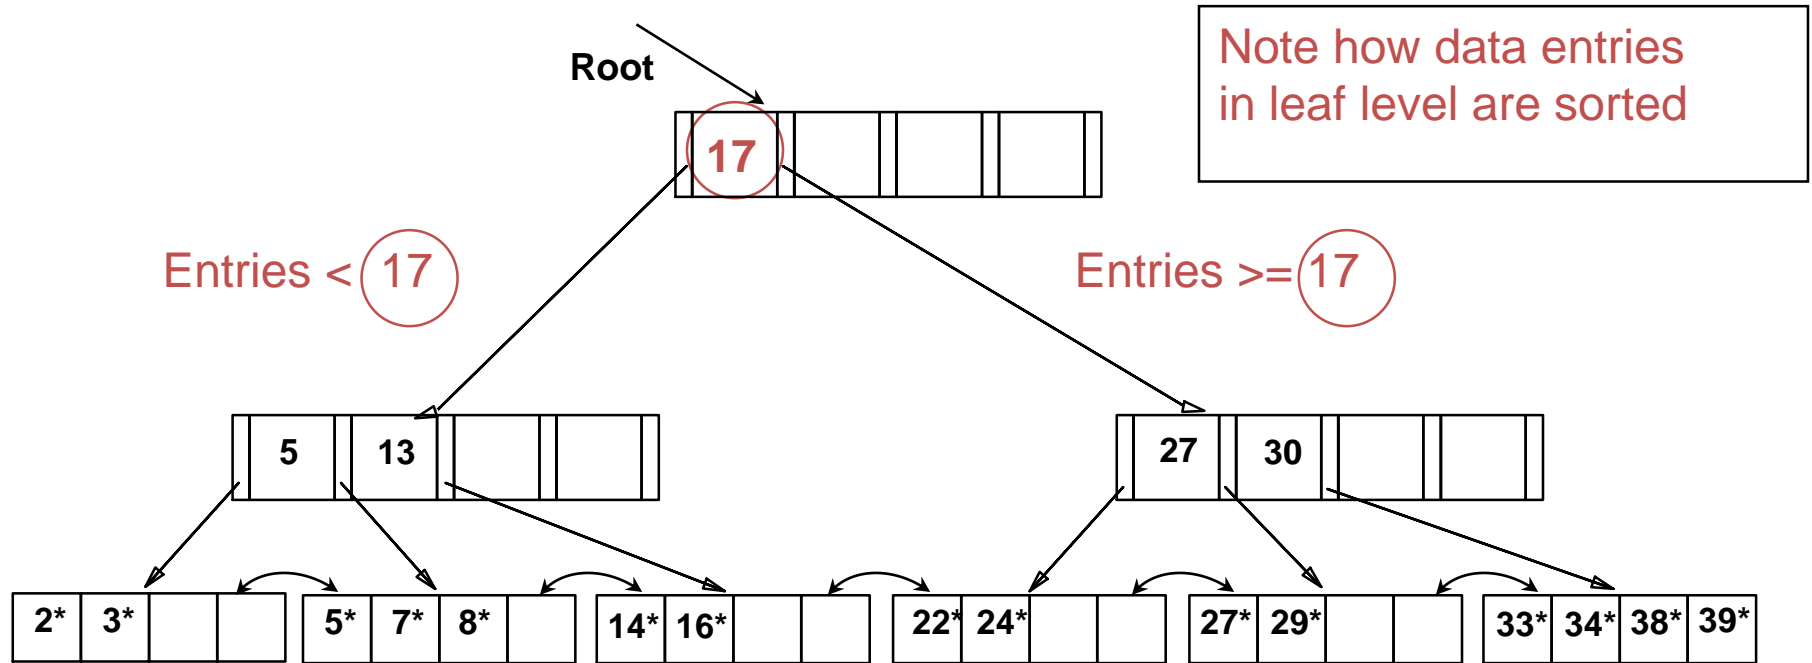
# RDF3X 6 indexes (Hexastore)









# B+ Tree



# RDF3X Query Processing

`SELECT ?x where { alice, lives_in, ?x }`

- Lookup ids :       **alice** → 1, **lives\_in** → 2
- Read results while prefix (1,2) matches: **(1,2,4)**




# RDF3X Query Processing

SELECT ?x where { Einstein, invented, ?x }

- Lookup ids **Einstein** → 16, **invented**
- Read prefix matches (16,24): **(16,24,567)** **(16,24,876)**

already  
sorted

All triples in  
(s,p,o) order









# RDF3X Storage and Indexing

Build clustered indexes for all **six permutations**

- SPO, POS, OSP to cover *all possible* triple patterns
- SOP, OPS, PSO to have *all sort orders* for patterns with two variables

**Triple table no longer needed, all triples in each index**

# How do the indexes work together ?




# Now, how join two **(2)** triple patterns ?

```
SELECT ?x where {  
    ?x, lives_in, Paris.  
    ?x, works_for, EDF  
}
```

- Naïve-way : evaluate one triple pattern at-a-time and then join (=intersect) the results

# Recall Dictionnary



dictionary

id	string
1	alice
2	works_for
3	lives_in
4	EDF
5	Paris


Giant-Table

subject	pred.	object
1	2	4
1	3	5




# RDF3X : Sort-Merge-join

- For example, we decide to use POS & OPS index for 1<sup>st</sup> & 2<sup>nd</sup> pattern, respectively.
  - we will see next why we did this choice




# RDF3X : Sort-Merge-join

- Scan both inputs: join matching values OR skip
- Idea : advance pointer with lower value




# RDF3X : Sort-Merge-join

- We access POS[3.5.**1**] and OPS[2.4.**1**]
  - we find **1** on both sides -> query result




# RDF3X : Sort-Merge-join

- We access POS [3.5.**7**]
  - then we know that OPS[2.4.{**2**..**6**}] are not results



# RDF3X : Sort-Merge-join

- We access OPS[2.4.**14**] then we know that PSO[3.5.{**6**..**13**}] are not results




# For unsorted triples ? Nested Loop

for each element ***x*** of **A**

for each element ***y*** of **B**

compare ***x*** with ***y***

## Nested Loop Join



# RDF3X Query Evaluation

- How to join n-triple patterns  $t_1 \dots t_n$ ?

```
SELECT ?x where {  
  ?x, lives_in, ?y . t1  
  ?x, works_for, ?z  
  . t2  
  ?y, isLocatedIn, "US" . t3  
  ?z, isLocatedIn, "US" . t4  
}
```


Classical relational problem: choose a left-deep join order, like

$((t_1 \text{ Join } t_2) \text{ Join } t_4) \text{ Join } t_3$

# RDF3X Query Evaluation

Classical problem: chose a left-deep plan, like

$((t_1 \text{ Join } t_2) \text{ Join } t_4) \text{ Join } t_3$





# Scenario 1 : no triples using property

`lives_in`

- then

$( (t_1 \text{ Join } t_2) \text{ Join } t_4) \text{ Join } t_3$

is optimal as it immediately discovers that the query has no answer

```
SELECT ?x where {  
    ?x, lives_in, ?y . t1  
    ?x, works_for, ?z . t2  
    ?y, isLocatedIn, "US" . t3  
    ?z, isLocatedIn, "US" . t4  
}
```

# Scenario 2: no triples using property `isLocatedIn`

- then

$( (t_1 \text{ Join } t_2) \text{ Join } t_4) \text{ Join } t_3$

is **not** optimal as it computes  $(t_1 \text{ Join } t_2)$  before understanding that the query is empty

```
SELECT ?x where {
```

```
    ?x, lives_in, ?y .           t1  
    ?x, works_for, ?z .         t2  
    ?y, isLocatedIn, "US" .     t3  
    ?z, isLocatedIn, "US" .     t4  
}
```

**Scenario 3 : `isLocatedIn` is very rare and `lives_in` is more frequent than `works_for`**

- then

$(t_4 \text{ Join } t_2) \text{ Join } (t_3 \text{ Join } t_1)$

is optimal (in average) as it is likely to keep intermediate results low (but this plan is not left deep!!!)

```
SELECT ?x where {
```

```
    ?x, lives_in, ?y      .      t1
    ?x, works_for, ?z     .      t2
    ?y, isLocatedIn, "US" .      t3
    ?z, isLocatedIn, "US"   t4
}
```

**COLUMN-STORES**

# Column-store

Relational database, storing relations as columns.

(C-Store, MonetDB and VectorWise, Ingres, IBM&MS Analytics)

product	country	sales
car	US	40K
bike	US	7K
	...	



**col1**

car@bike

**col2**

US@US


**col3**

40K@7K

..

# Today, any Relational Datawarehouse is a Column-Store


- «*Do not call them column-stores, just modern DB systems.*»  
Stratos Idreos



- Tables in a DW can have more than ten dimensions

# Today, any Relational Datawarehouse is a Column-Store

- «*Do not call them column-stores, just modern DB systems.*»  
Stratos Idreos



- Analytical query just uses a few of them

# Rows become crowded

**row1**

id1@customer1@  
region1@sales\_  
rep1@discount1  
@product1  
@weather1@date  
1@state1...

**row2**


id2@customer2@  
region2@sales\_  
rep2@discount2  
@product2  
@weather2@date  
2@state2

...






# Worst case : visit whole row



<b>row1</b>	<b>id1</b> @ <b>customer1</b> @ region1@sales_ rep1@discount1 @ <b>product1</b> @weather1@ <b>date</b> <b>1</b> @state1...	<b>row2</b>	<b>id2</b> @ <b>customer2</b> @ region2@sales_ rep2@discount2 @ <b>product2</b> @weather2@ <b>date</b> <b>2</b> @state2	...
-------------	---	-------------	--	-----

- **Not only a reading-problem : entire sets of rows have to be loaded into memory from disk!**
  - This wastes bandwidth



# Rows are stored in pages


**row1** **id1**@**customer1**@  
region1@sales\_  
rep1@discount1  
@**product1**  
@weather1@**date**  
**1**@state1...

**row2** **id2**@**customer2**@  
region2@sales\_  
rep2@discount2  
@**product2**  
@weather2@**date**  
**2**@state2

...

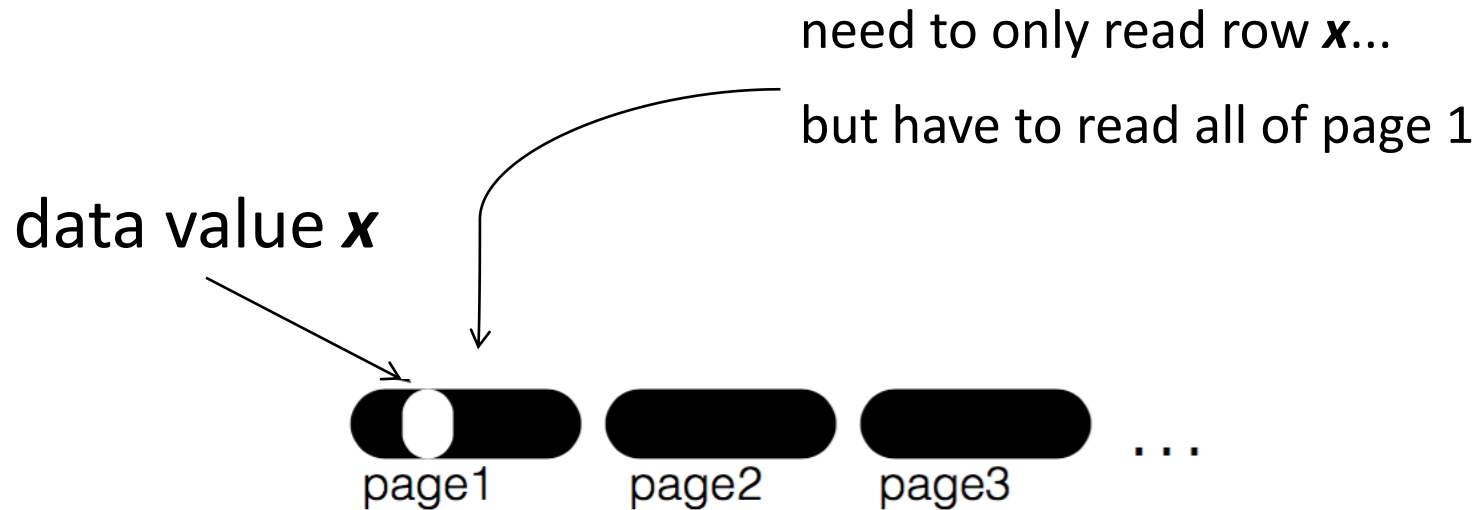
page

# Pages are stored within files




file

# Constraint : Pages are read as a whole!



# Better to « pay as you go »!



<b>col11</b>	<b>id1@id2</b>	<b>col12</b>	<b>customer1@customer2</b>	<b>col13</b>	region1@region2
<b>col14</b>	sales_rep1@sales_rep2	<b>col15</b>	discount1@discount2	<b>col16</b>	<b>product1@product12</b>
<b>col17</b>	weather1@weather2	<b>col18</b>	<b>date1@date2</b>	<b>col19</b>	state1@state2

# Store columns on separate pages!



<b>col1</b>	<b>id1@id2</b>	<b>col2</b>	<b>customer1@customer2</b>	<b>col3</b>	region1@region2
<b>col4</b>	sales_rep1@sales_rep2	<b>col5</b>	discount1@discount2	<b>col6</b>	<b>product1@product12</b>
<b>col7</b>	weather1@weather2	<b>col8</b>	<b>date1@date2</b>	<b>col9</b>	state1@state2

**col1** id1@id2

**col2** custome  
r1@cust  
omer2

**col3** region  
1@regi  
on2

**col4** sales\_r  
ep1@sal  
es\_rep2

**col5** discount  
1@discou  
nt2@

**col6** produc  
t@1pro  
duct12

**col7** weather  
1@weath  
er2

**col8** date1@  
date2

**col9** state1  
@state  
2

file



CPU

registers

on chip cache

on board cache

memory

disk

data move

col12

custome  
r1@cust  
omer2



# Column Stores

The state of the art solution for

1. analytical
2. read-mostly queries
3. on wide-relations
4. supporting only batch-updates

# But wait...

- This does not mean that columnstores will automatically work well for RDF data..
  - for example, RDF property-tables are not wide, cluster-tables are wide however
- Let's see...

# Logical Model for RDF ColumnStores : Property-Tables

- A relational table for each single RDF property.

Giant-Table	subject	predicate	object
	alice	works_for	EDF
	alice	lives_in	Paris
	...		

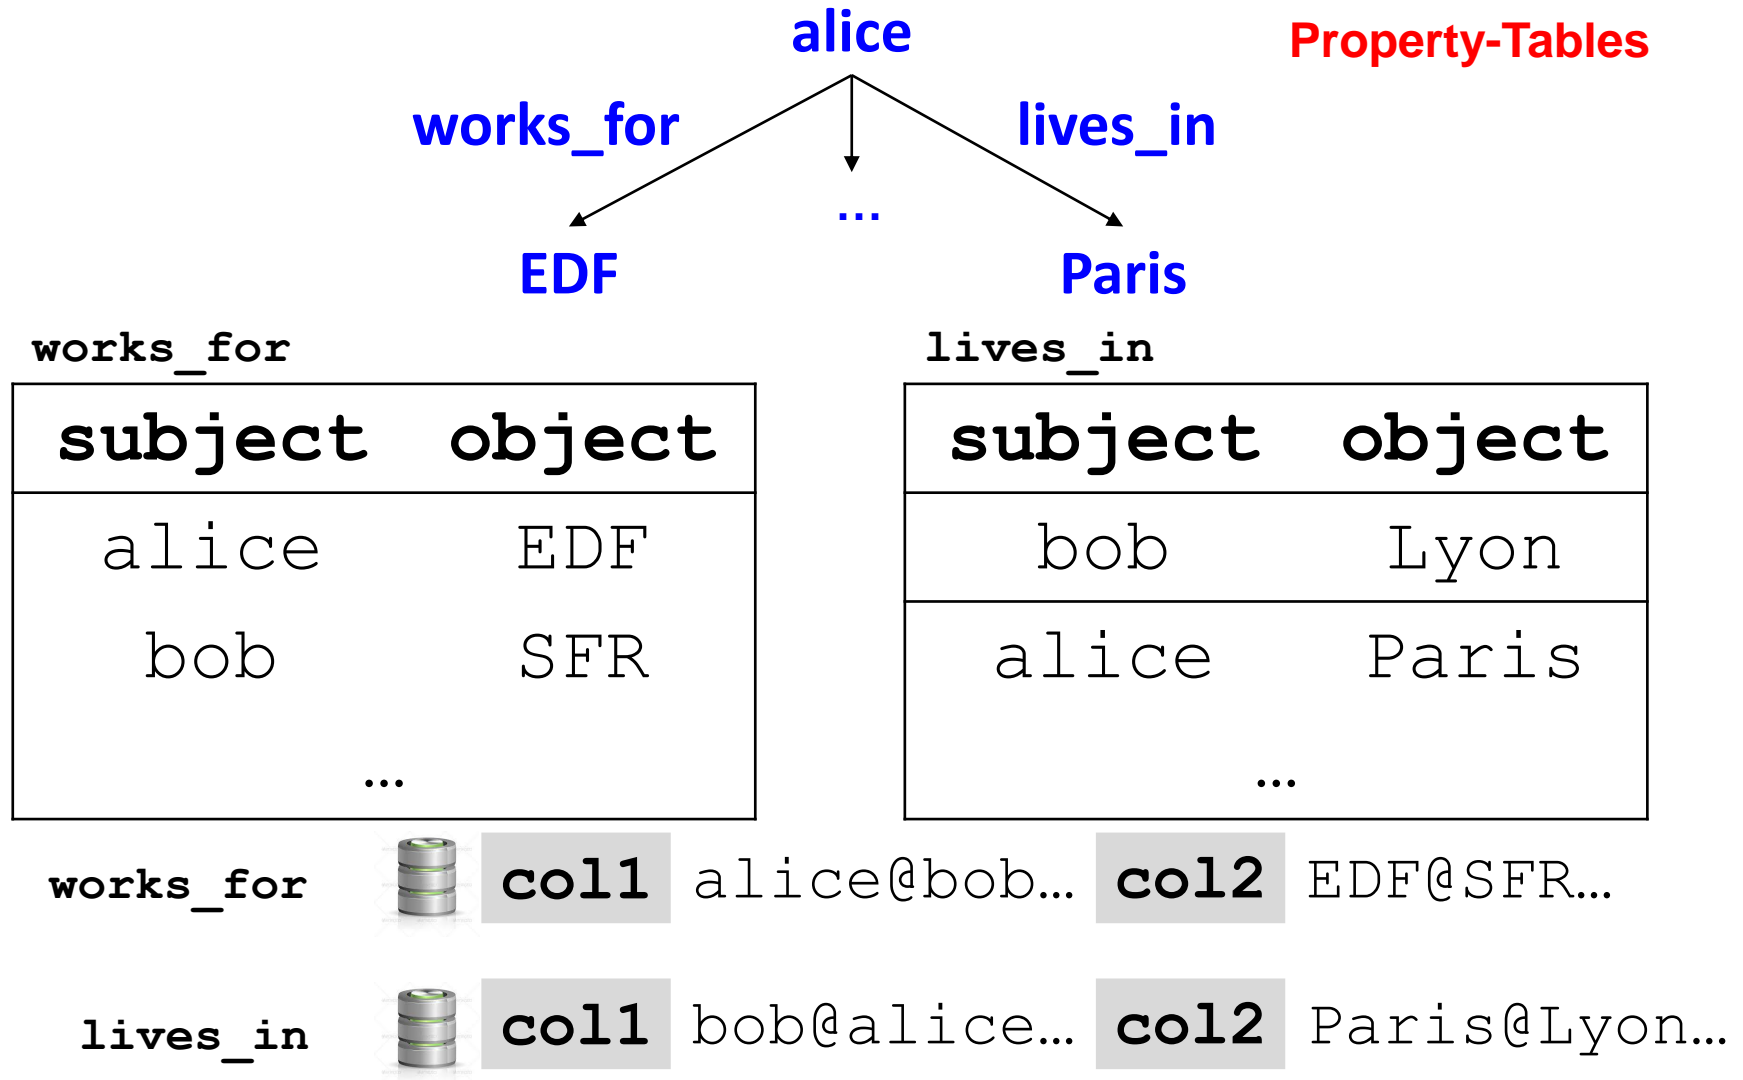
**works\_for**

subject	object
alice	EDF
...	

**lives\_in**

subject	object
alice	Paris
...	


# RDF in a Column-store (MonetDB, C-Store, Virtuoso)



# Neat advantage : column-projection

- Fast queries if only subject or object of a triple are accessed, not both

```
SQL> SELECT subject FROM lives_in
```

`lives_in`  **col1** alice@... **col2** Paris@...

- All results can be found in the first column

# Advantages of using column stores

- Allows for a very compact representation
- Exploits merge joins
- A column contains “homogeneous” values, where many values repeat and thus compression is more effective



**col1**

car@bike

**col2**

US@US

**col3**

40K@7K

..

# Disadvantages of using column stores

- Need to recombine columns if subject and object are accessed



**col1** car@bike **col2** US@US **col3** 40K@7K ..

- Inefficient for triple patterns with predicate variable  
<alice ?p bob>
- Big question : when to reconstruct a tuple ?  
(alap)

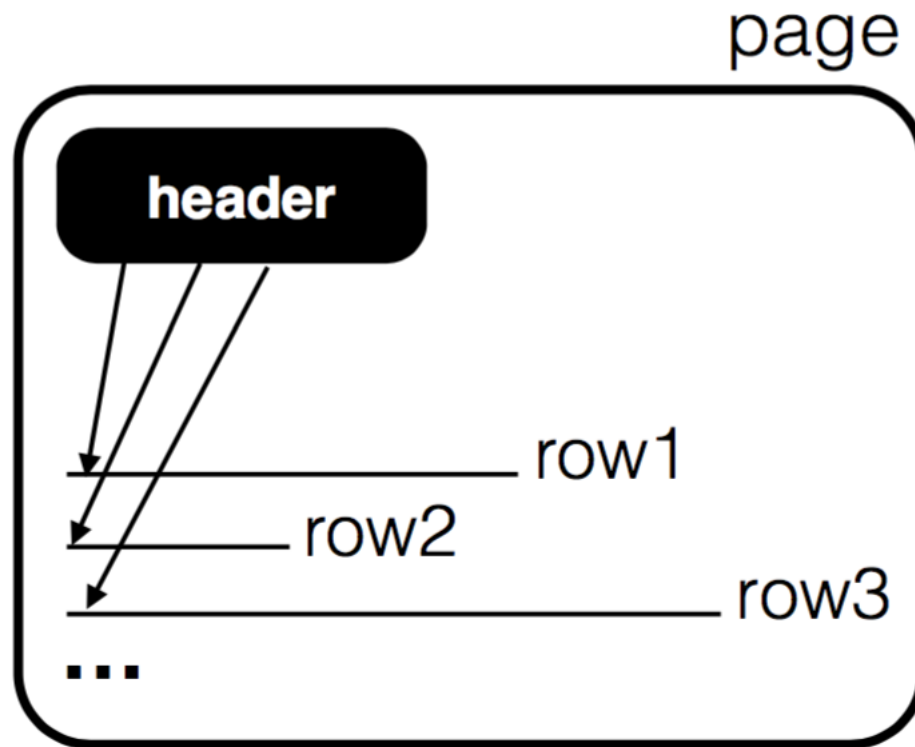


# Advantages of Column Stores


1. Tuple Headers Stored Separately
2. Optimizations for fixed-length tuples
3. Column-oriented data compression
4. Carefully optimized merge-join code

**TUPLE HEADERS**

# Page headers in Rowstores



# Tuple headers in Rowstores



| tuple header | >> | row<sub>of size 2</sub> |

# Tuple headers

- Metadata **at the beginning of the row**
  - Insert transaction timestamp
  - number of attributes in tuple (useless for RDF)
  - NULL flags
- Postgres: 27 byte tuple header + 8 bytes two-column tables

| header |      >>      | row<sub>of size 2</sub> |

# Postgres Tuple Header

Field	Type	Length	Description
t_xmin	TransactionId	4 bytes	insert XID stamp
t_xmax	TransactionId	4 bytes	delete XID stamp
t_cid	CommandId	4 bytes	insert and/or delete CID stamp (overlays with t_xvac)
t_xvac	TransactionId	4 bytes	XID for VACUUM operation moving a row version
t_ctid	ItemPointerData	6 bytes	current TID of this or newer row version
t_infomask2	int16	2 bytes	number of attributes, plus various flag bits
t_infomask	uint16	2 bytes	various flag bits
t_hoff	uint8	1 byte	offset to user data
null bitmap		optional	
object id		optional	

# Postgres Tuple Header

- Some information on visibility of a tuple for current transaction snapshot or newer version (needed for snapshot isolation algorithm)
  - t\_xmin TransactionId 4 bytes insert XID stamp
  - t\_xmax TransactionId 4 bytes delete XID stamp
  - t\_cid CommandId 4 bytes insert CID stamp (actual a UNION struct)
  - t\_ctid ItemPointerData 6 bytes current TID of this or newer row version
- How long is this row? Is it variable length? Does it have NULLs?
  - t\_natts int16 2 bytes number of attributes
  - t\_infomask uint16 2 bytes various flag bits
    - e.g. HAS\_NULL | HASVARWIDTH | HASOID | locks(!)
- t\_hoff uint8 1 byte /\* sizeof header incl. bitmap, padding \*/
- null bitmap (optional)
- object id(optional)

# Tuple headers in Columnstores

- Puts header information in separate columns and can selectively ignore it
  - #attributes is always two
  - in some cases, NULL values are totally avoided
- Column-store effective tuple width : ~8 bytes
- Reading a tuple takes 4-5 times less time than Postgres (27 + 8bytes), so does a simple table scan.




# Rowstore vs Columnstore Headers

**Postgres Header (27 bytes)**




**Tuple Data**

**Column Store Header**



**Tuple Data**



# **OPTIMIZATIONS FOR FIXED-LENGTH TUPLES**

# Optimizations for variable-length tuples

- 1 attribute **variable length** => the whole tuple is
- Common case: row-stores designed for this

```
row1 idcustomer1@firstname1@lastname1@...
```

- Tuples located/iterated via pointers in the page header (instead of address offset calculation)

# Optimizations for fixed-length tuples

- In columnstores, with a dictionary encoding, fixed length val. are stored/accessed as **arrays**

<b>A</b>	a0
	a1
	a2
	a3
	..

Positional lookup

$$a(i) = A + i * width(A)$$

# Optimizations for fixed-length tuples

- Every RDF property table has two columns

works_for	
subject	object
alice	EDF
...	

- Store each on disk in a page (eg, 64K)
  - in SW-Store, they found this suboptimal for queries accessing both values

# Hybrid Storage (SW-Store)

- A page slot contains data from both columns


**P**

s0	o0
s1	o1
s2	o2
s3	o3
...	...

$$o(i) = P + i * (\text{width}(S) + \text{width}(O)) + \text{width}(S)$$

# +Index on hybrid Column (SW-Store)

- clustered B+ tree index on subject
- unclustered B+ tree index on object.



## Clustered/unclustered

- Clustered = records close in index are close in data
- Unclustered = records close in index may be far in data

# More optimizations : Id-Table

- Maintain a **single-column** table that contains all triple subjects (much better if ordered)

Id-Table

$id_{\text{Obama}}$
$id_{\text{Alice}}$
$id_{\text{Paris}}$
• •



# Id-Table


- If property **P** is NOT multivalued (eg. birthday) we can avoid to store subject id

Id-Table	$id_{Obama}$
	$id_{Alice}$
	$id_{Paris}$
	• •

<b>P</b>	o0
	o1
	o2
	o3
	• •


# Id-Table

- If property **P** is NOT multivalued (eg. birthday) we can avoid to store subject id



# Nulls are back!

- If property **P** is NOT multivalued (eg. birthday) we can avoid to store subject id



# Id-Table : Dealing with Nulls

- Goal : remove nulls from the column
- There is not a single optimal solution
- The point is to find a tradeoff between the **size** and **manegeability** of the compressed column
  - this depends only on the **data-distribution**

# It's matter of « density »

D. Abadi [Column-Stores For Wide and Sparse Data](#)

Dense



Not dense nor sparse



Sparse



# Case 1: «dense» data and long null sequences

**Title**

t1
t2
t3
t4
NUL L

**Language**

NUL L
FR
EN
NUL L
NUL

# Case 1: Low overhead for «dense» data

- Store indexes of non-null values

**Title**

t1
t2
t3
t4
NUL L

**Title**

t1
t2
t3
t4

**Range [1-4]**

# Case 1: Low overhead for «dense» data

- Store indexes of non-null elements

Language

NUL L
FR
EN
NUL L
NUL

Language

FR
EN

Range [2-3]



## Case 1: Low overhead for «dense» data

- Store indexes of non-null elements
- If data is dense, then

$$|\text{range information}| \ll |\text{nulls}|$$

## Case 2: data not «dense» nor «sparse»

- Store a bitmap index (0=NULL)

Copyright

2001
NUL L
1985
NUL L
1995

## Case 2: data not «dense» nor «sparse»

- Store a bitmap index (0=NULL)

Copyright Bit:101011

2001
NUL L
1985
NUL L
1995

## Case 2: data not «dense» nor «sparse»

- Store a bitmap index (0=NULL)

Copyright Bit:101011

- Overhead = 1bit per value

2001
NUL L
1985
NUL L
1995

# Case 3 : sparse data

- Store the list of non-null ids

**Author**

Hugo
NUL L
NUL L
NUL L

**Artist**

NUL L
Dylan
NUL L
NUL L

# Case 3 : sparse data

- Store the list of non-null ids

**Author**

Hugo
NUL L
NUL L
NUL L

**Author**

List:1

Hugo
------

# Case 3 : sparse data

- Store the list of non-null ids

**Artist**

NUL L
Dylan
NUL L
NUL L

**Artist**

List:2

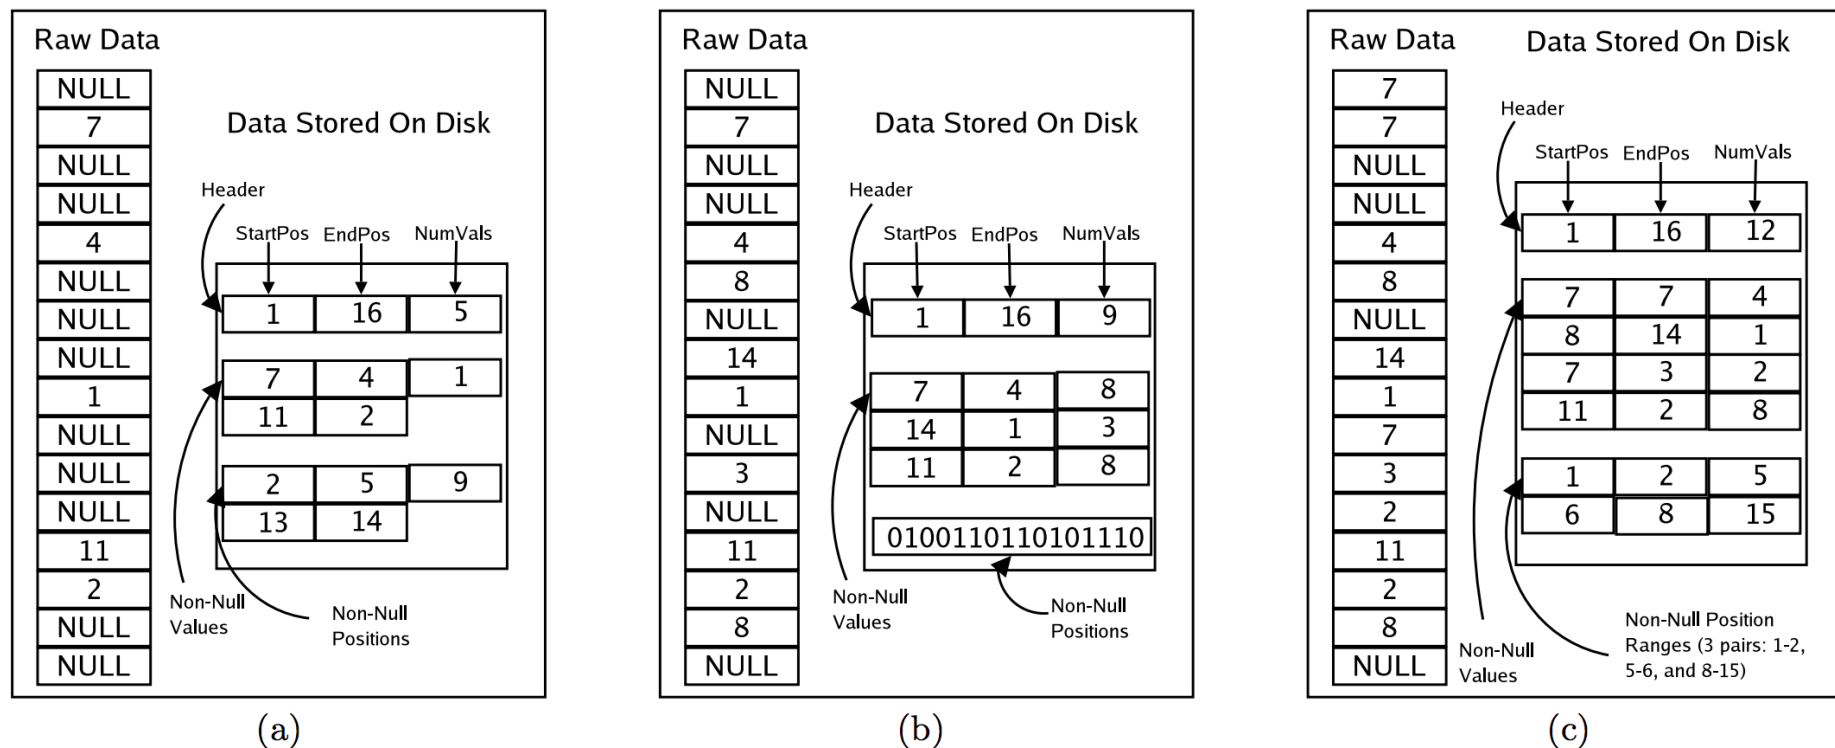
Dylan

## Case 3 : sparse data

- Store the list of non-null ids
- If data is sparse

$$|\text{List}| \ll |\text{nulls}|$$





**Figure 1: Positions represented using a list (a), a bit-string (b), and as ranges (c) for sparse columns**

# Back to Query Evaluation

SELECT ?y ?z where {

    ?x, author, Hugo

. t<sub>1</sub> **10<sup>-6</sup>**

    ?x, title, ?y

. t<sub>2</sub> **2\*10<sup>-2</sup>**

    ?x copyright ?z


. t<sub>3</sub> **10<sup>-3</sup>**

}

triple  
pattern  
selectivity


(t<sub>1</sub> Join t<sub>3</sub>) Join t<sub>2</sub>

# Plan



**Author**

List:1

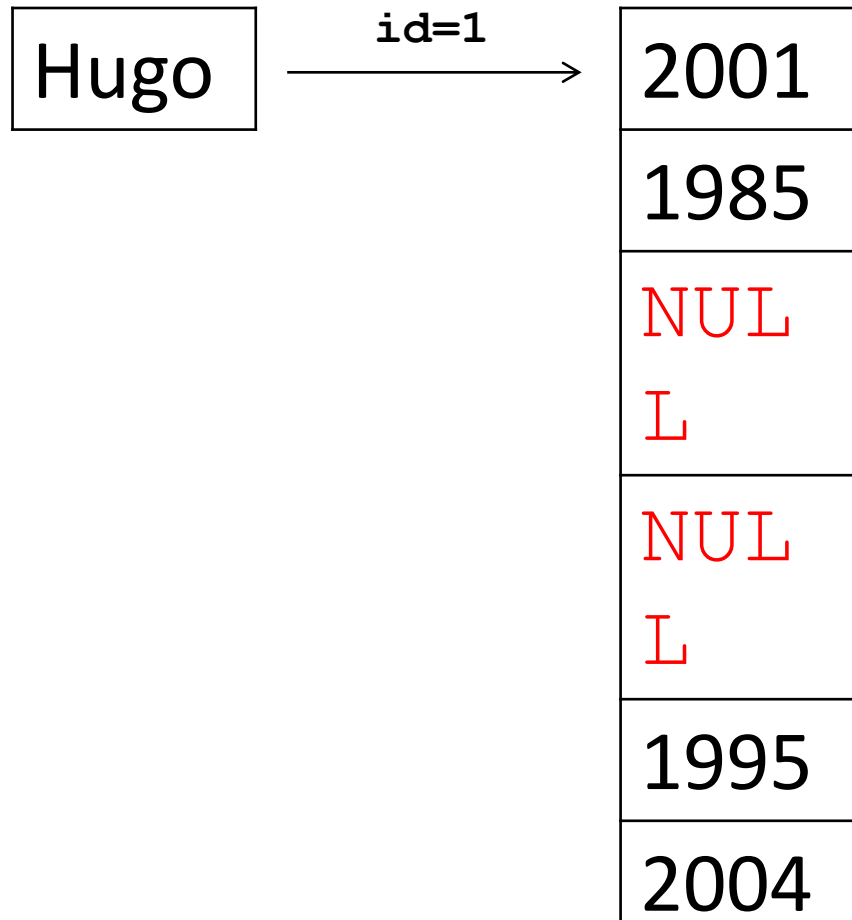


**Author**

List:1

**Copyright**

Bit:110011



**Author**

List:1

Hugo

**Copyright**



Bit:110011

2001
1985
NUL L
NUL L
1995
2004


**Title**

Range[1-5]

t1
t2
t3
t4



# Query



```
SELECT ?y ?z where {
```

```
    ?x, author, Hugo
```

```
    . t1 10-6
```

```
    ?x, title, ?y
```

```
    . t2 2*10-2
```


```
    ?x copyright ?z
```

```
    . t3 10-3
```

```
}
```

$(t_2 \text{ Join } t_3) \text{ Join } t_1$

# Plan





**Title**  
Range[1-7]

t1
t2
t3
t4
t5
t6
t7

**Copyright**  
Bit:110011

2001
1985
NUL L
NUL L
1995
2004

**Author**  
List:1

Hugo
------

↙ ↘  
id=  
1,2,3,  
4,5,6,  
7,...

↙ ↘  
id=  
1,~~2,3,~~  
4,5,6,  
7,...

↘  
id=1

# **COLUMN-ORIENTED DATA COMPRESSION**

# Column-oriented data compression

- Since each attribute is stored separately (even within a slot), it can be compressed separately using the best algorithm
  - for example, the subject ID column, a monotonically increasing array of integers
  - data from the same domain tend to show locality
- It is often possible also to operate directly on compressed representations
  - Bandwidth requirements are reduced when transferring compressed data

**SO FAR SO GOOD. NOW WHO  
WINS?**


# There is no single winner

	<i>RDF-3x</i>	<i>VOS [6.1]</i>	<i>VOS [7.1]</i>	<i>MonetDB</i>	<i>4Store</i>
% of queries for which tested system is fastest	20.9%	0.0%	22.6%	56.5%	0.0%
Total workload execution time (hours)	27.1	20.9	20.8	38.6	72.2
Mean (per query) execution time (seconds)	7.8	6.0	6.0	11.1	20.8

Table 1.1: Summary of results over WatDiv 100M RDF triples, 12500 SPARQL queries.

# Mini-Projet : Moteur RDF (Partie 1)

- Implémenter un moteur de requêtes pour données RDF utilisant l'une des approches vues en cours :
  - Hexastore, Columnstore, Graphstore



# Mini-Projet : Moteur RDF (Partie 2)

- Évaluer les performances du système réalisé
  - comparer avec 1 autre système + Jena
- Fournir un système fonctionnel à ses collègues (dans les délais) fait partie de l'évaluation